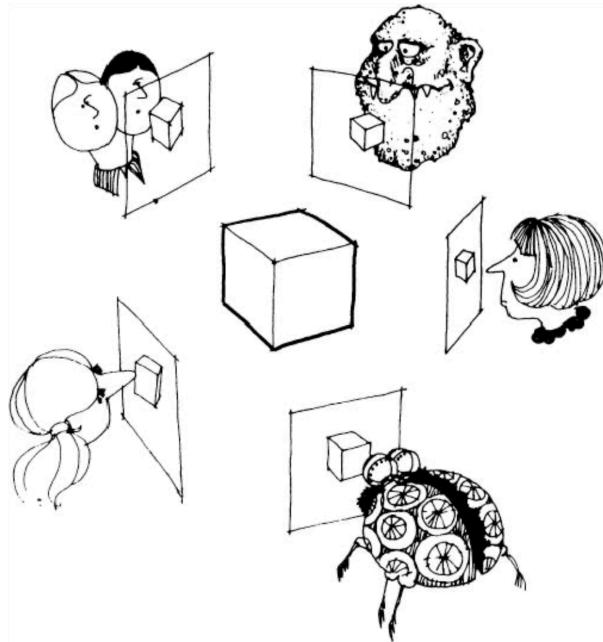


PatchMatch in Multi-View Stereo

Yiming Xie

2020.6.21



Many slides adapted from E. Dunn, S. Shen, Y. Furukawa, M. Pollefeys, and others

ETH 3D Benchmark(High Res.)

	Method	all	high-res multi-view	indoor	outdoor
1	DeepPCF-MVS	80.84		88.10	88.56
2	DeepC-MVS_fast	79.62		86.82	86.47
3	DeepC-MVS	79.50		86.80	86.53
4	3Dnovator+	77.84		85.48	84.47
5	MG-MVS			83.41	83.45
6	3Dnovator	76.31		83.38	82.31
7	CLD-MVS			82.31	81.65
8	AP-MVS			82.00	81.11
9	MAR-MVS			81.84	80.70
10	ACMP	74.13		81.51	80.57
11	ACMM	73.20		80.78	79.84
12	AdaColmap			80.58	79.42
13	PCF-MVS	73.52		80.38	78.84
14	OpenMVS	72.83		79.77	78.33
15	TAPA-MVS	73.13		79.15	77.94
16	PLC	70.83		78.05	76.37
17	LS3D			76.95	74.82
18	F-COLMAP			76.38	74.33
19	LTVRE_ROB	69.57		76.25	74.54
20	ACMH+	68.96		76.01	74.01
21	ACMH	67.68		75.89	73.93
22	MSDG			73.36	70.99
23	COLMAP_ROB	66.92		73.01	70.41
24	OpenMVS_ROB	64.09		70.56	68.19
25	CMPMVS	51.72		70.19	68.16
26	LF4IMVS			64.02	62.19
27	PVSNet	57.27		61.67	59.27
28	PPMVS			53.68	51.64
29	ANet_high			50.57	46.10
30	Gipuma			45.18	41.86
31	PMVS	37.38		44.16	40.28
32	wuykxyi23d			39.76	33.18
33	PNet_i23d			38.26	35.42
34	MVE	26.22		30.37	25.89
35	wuyk23d			16.82	15.94
36	DeepMVS_CX				27.40
37	ITE_SJBPF				78.22



PatchMatch



Unknown



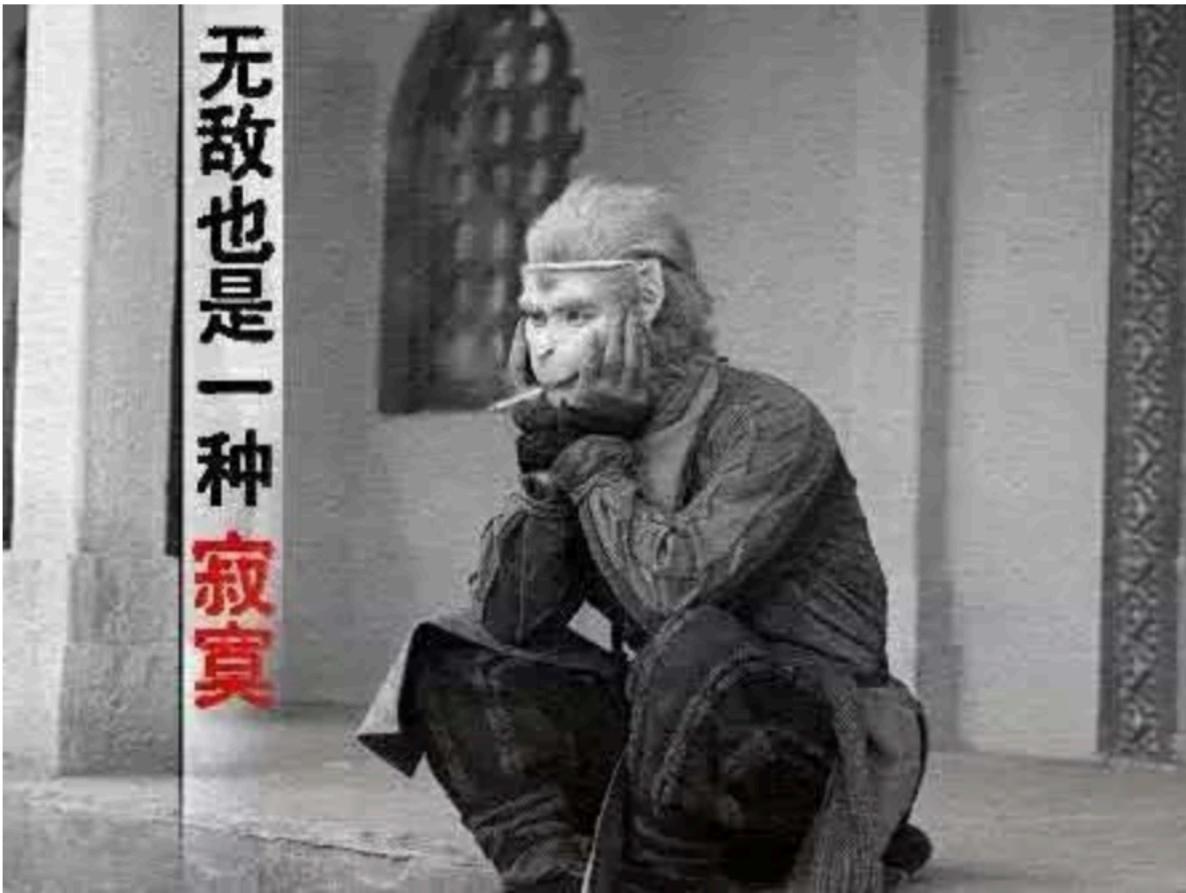
Others

2020.6.20



ETH 3D Benchmark(High Res.)

1	Method	all	high-res multi-view	indoor	outdoor
2	DeepPCF-MVS	80.84	88.10	88.56	86.73
3	DeepC-MVS_fast	79.62		86.82	86.47
4	DeepC-MVS	79.50		86.80	86.53
5	3Dnovator+	77.84		85.48	84.47
6	MG-MVS			83.41	83.45
7	3Dnovator				83.28
8	CLD-MVS				
9	AP-MVS				
10	MAR-MVS				
11	ACMP				
12	ACMM				
13	AdaColmap				
14	PCF-MVS				
15	OpenMVS				
16	TAPA-MVS				
17	PLC				
18	LS3D				
19	F-COLMAP				
20	LTVRE_ROB				
21	ACMH+				
22	ACMH				
23	MSDG				
24	COLMAP_ROB				
25	OpenMVS_ROB				
26	CMPMVS				
27	LF4IMVS				
28	PVSNet				
29	FFPMVS				
30	ANet_high				
31	Gipuma				
32	PMVS				
33	wuykxyi23d		39.76	33.18	59.51
34	PNet_i23d		38.26	35.42	46.79
35	MVE	26.22	30.37	25.89	43.81
36	wuyk23d		16.82	15.94	19.46
37	DeepMVS_CX				27.40
38	ITE_SJBPF				78.22



Topics

- **Introduction**
- **PatchMatch**
- **PatchMatch Stereo**
- **View Selection**

Introduction

Why Does it Matter?



UAV



Robotics



Augmented Reality

- **Goal: Sensing 3D Geometry**

Why Does it Matter?



UAV



Robotics

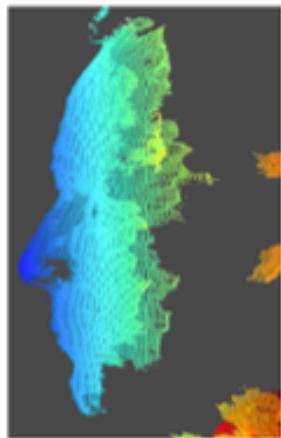


Augmented Reality

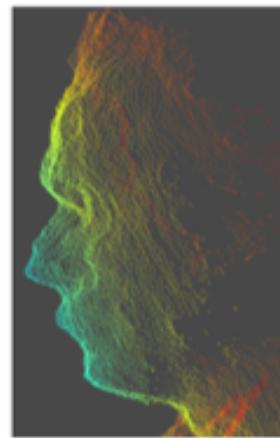
- Goal: Sensing 3D Geometry



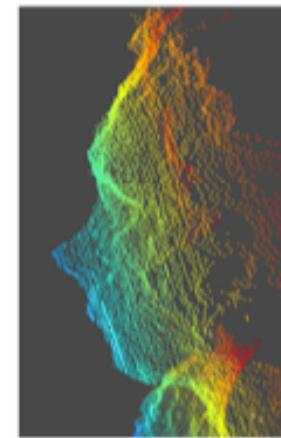
Depth Sensor



300mm



500mm



700mm



2500mm

Why Does it Matter?



UAV



Robotics



Augmented Reality

- **Goal: Sensing 3D Geometry**
- Among all, image-based methods provide a fast way of capturing accurate 3D content at a fraction of the cost of other approaches.

What is MVS

- Multi-view stereo (MVS): use stereo correspondence as their main cue and use more than two images to extract geometry from photographs.

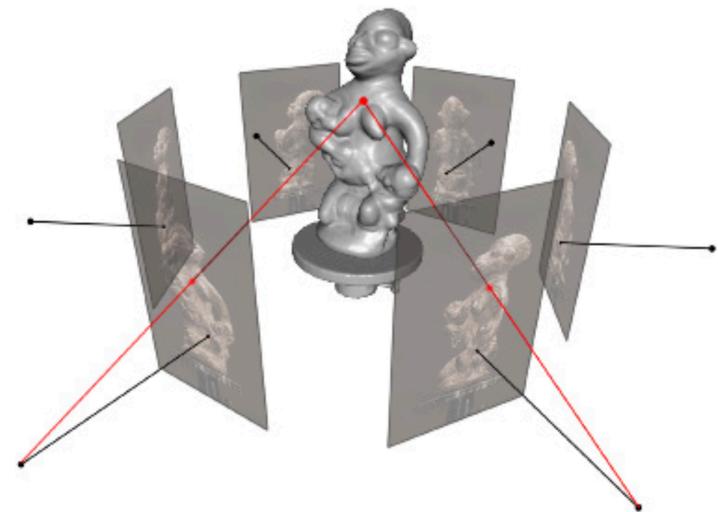
What is MVS

- Multi-view stereo (MVS): use stereo correspondence as their main cue and use more than two images to extract geometry from photographs.
 - Lambertian textured surfaces.
 - Known camera parameters.

What is MVS

- Multi-view stereo (MVS): use stereo correspondence as their main cue and use more than two images to extract geometry from photographs.
 - Lambertian textured surfaces.
 - Known camera parameters.

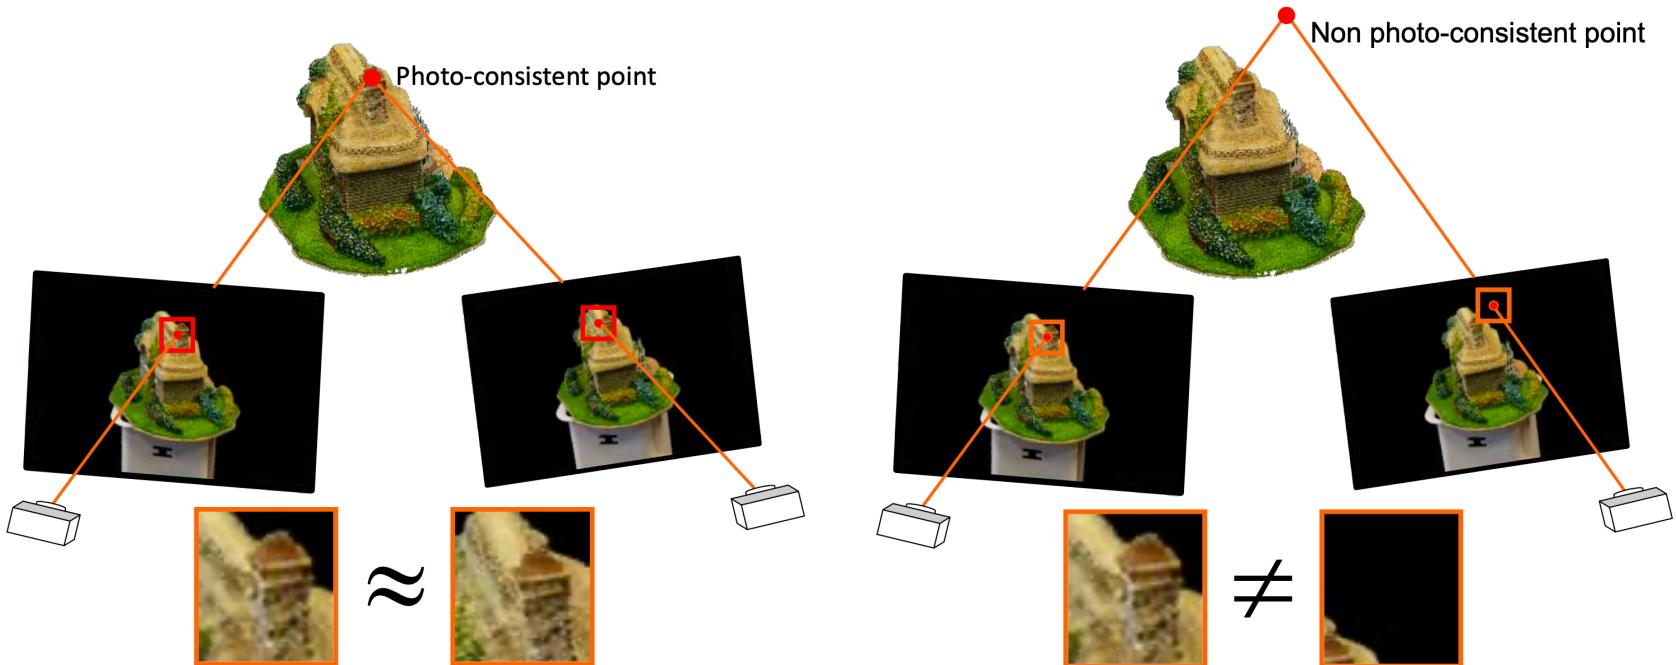
- Input: multiple images with calibrated cameras
- Output: dense 3d representation



Credit: Y. Furukawa

Multi-view stereo: Basic idea

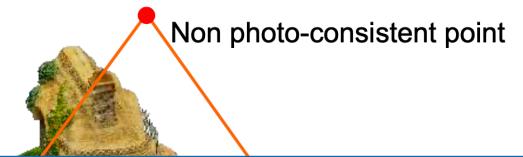
- Look for points in space that have photo-consistency.



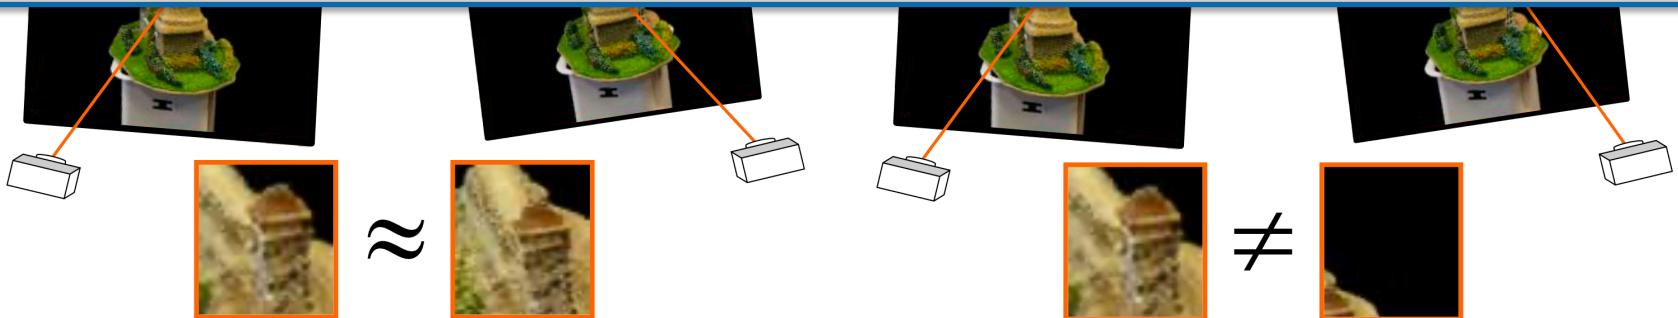
Credit: Shuhan Shen

Multi-view stereo: Basic idea

- Look for points in space that have photo-consistency.



Dense correspondence!!



Credit: Shuhan Shen

Summary

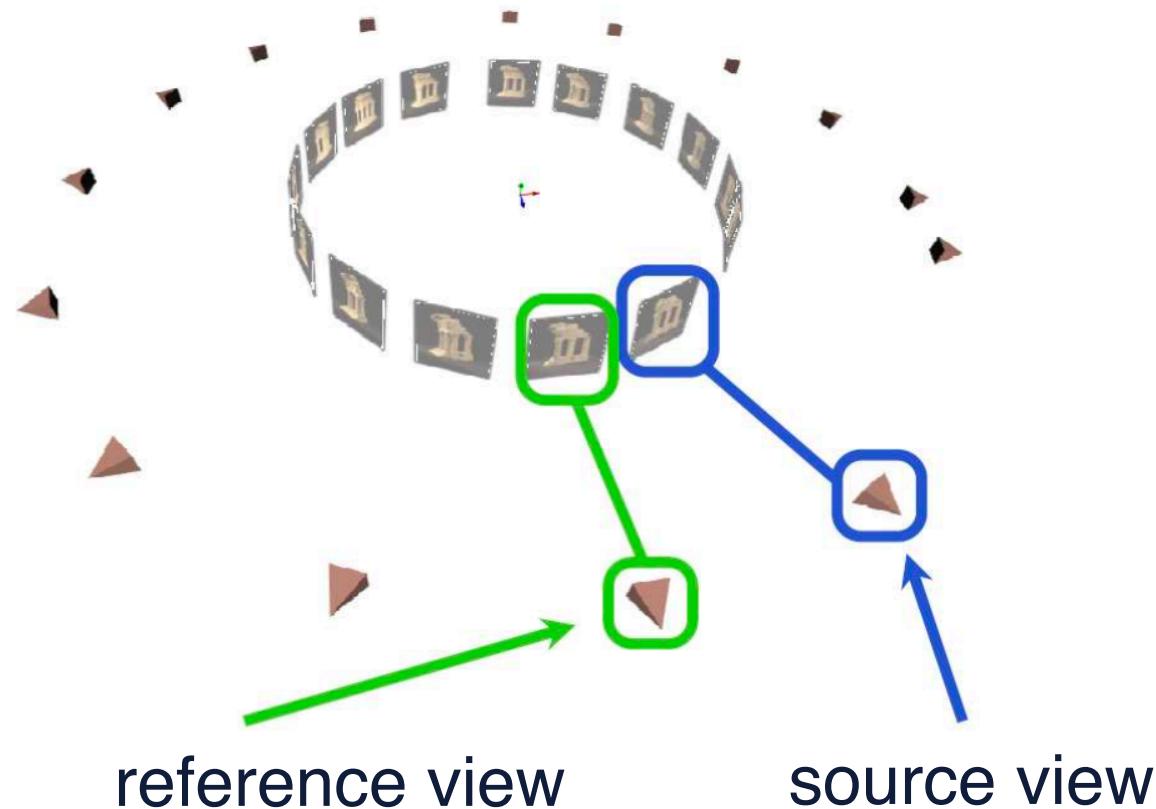
- **Why?**
- capture accurate 3D geometry, and image-based method is cheap.
- **What?**
- use stereo correspondence as their main cue and use more than two images to extract geometry from photographs.
- **How?**
- Look for points in space that have photo-consistency.

Depth-map Merging Based Approaches

Depth-map Merging Based Approaches

Depth-map Merging Based Approaches

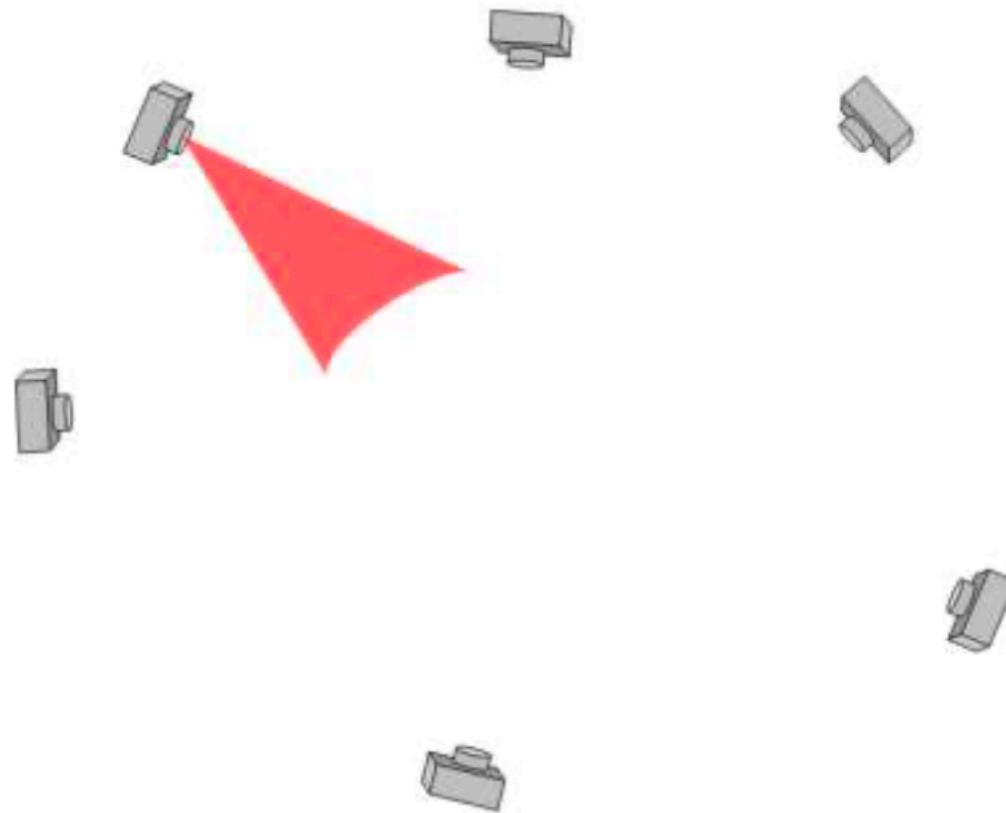
Step 1: Source view selection



Credit: Shuhan Shen

Depth-map Merging Based Approaches

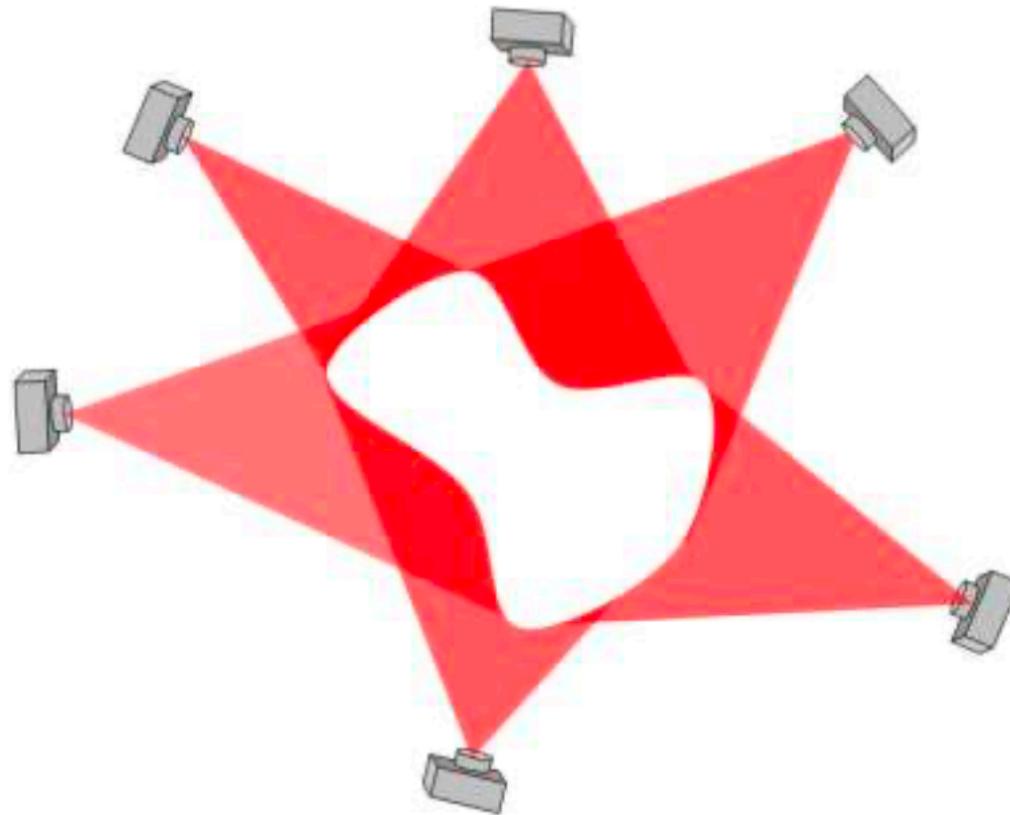
Step 2: Depth-map computation



Credit: Shuhan Shen

Depth-map Merging Based Approaches

Step 3: Depth-map merging



Credit: Shuhan Shen

- Step 1: Source view selection
- Step 2: Depth-map computation

- Step 3: Depth-map merging



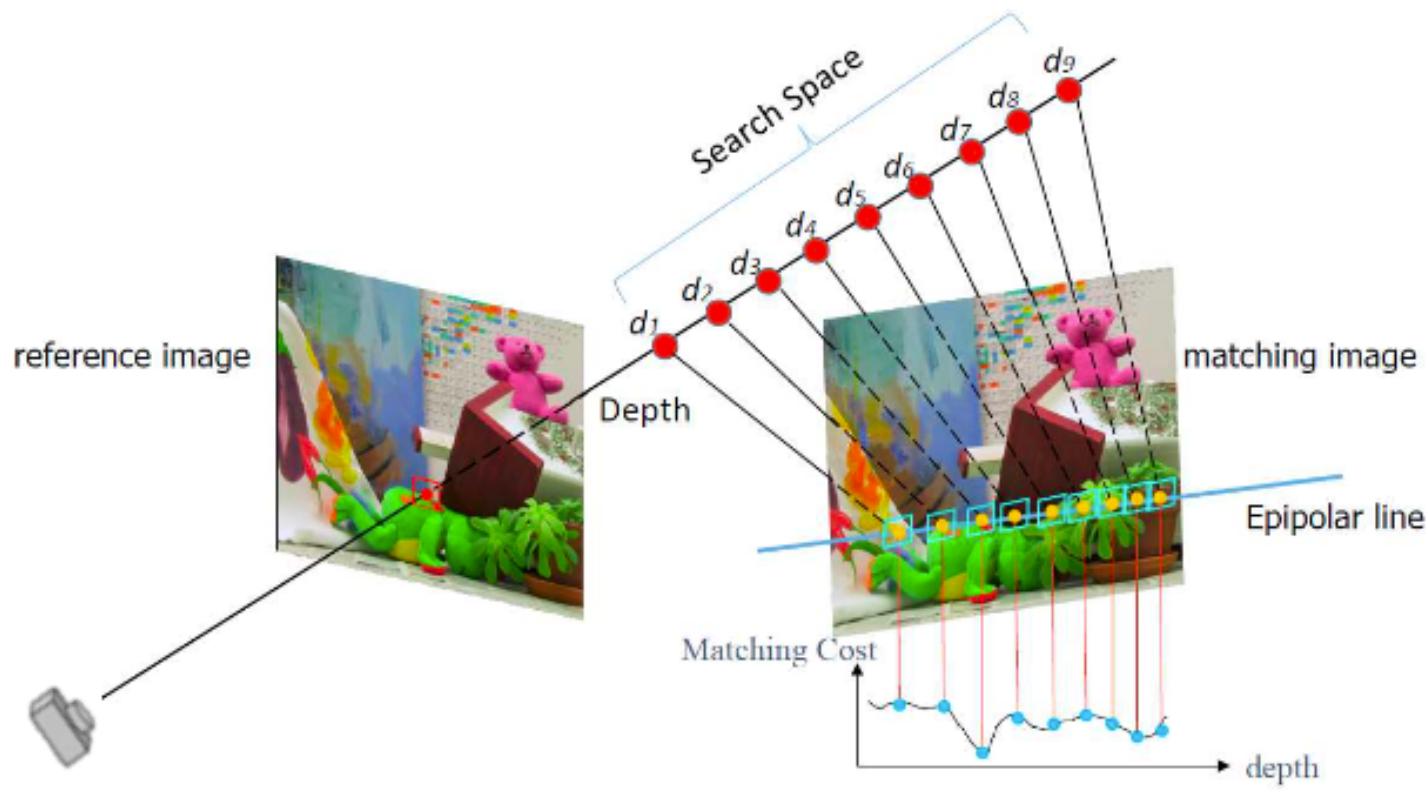
Key steps:

- 1.How to chose source images
- 2.How to compute depth map

How to compute depth map

Compute Depth Map: Basic idea

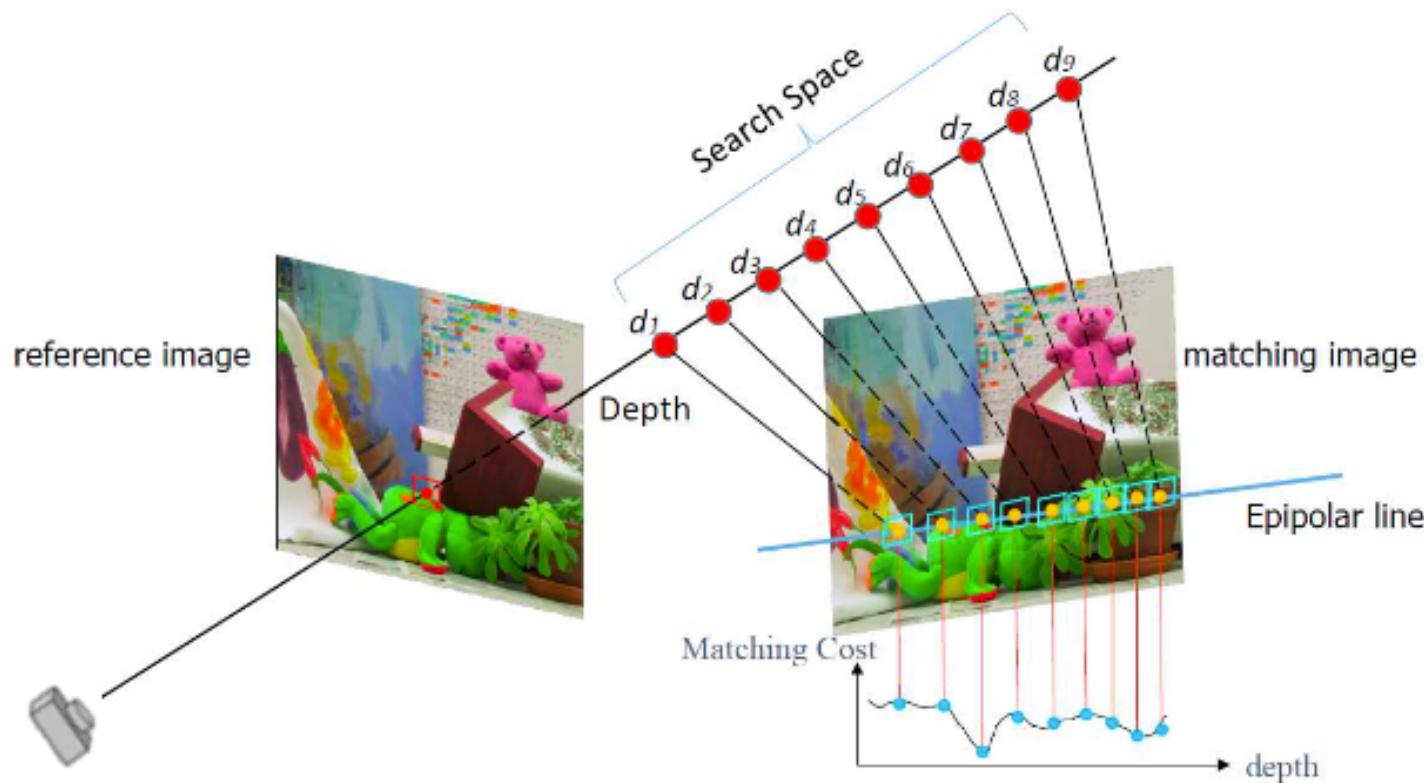
- Look for points in space that have photo-consistency.



Credit: E. Dunn

Compute Depth Map: Basic idea

- Look for points in space that have photo-consistency.



limit on high resolution images

Credit: E. Dunn

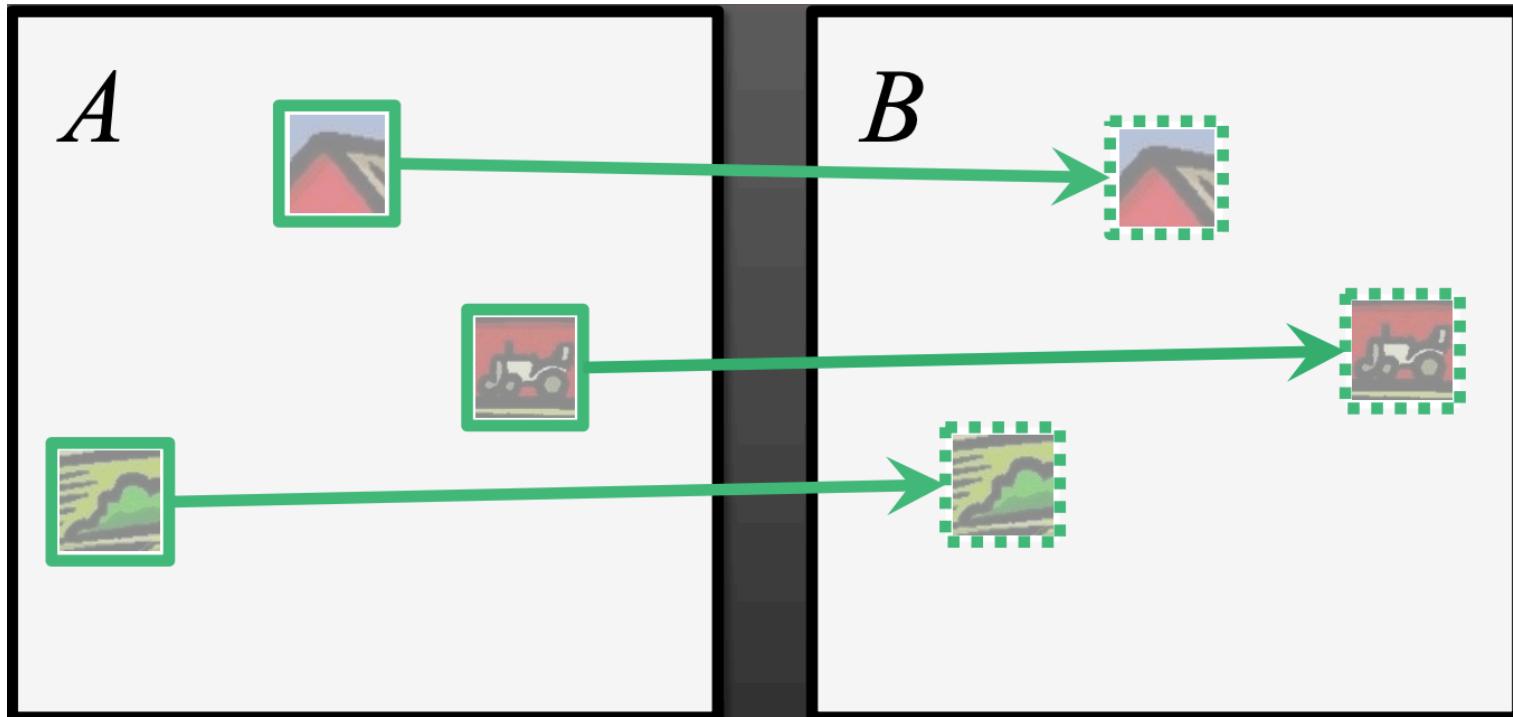
PatchMatch

PatchMatch

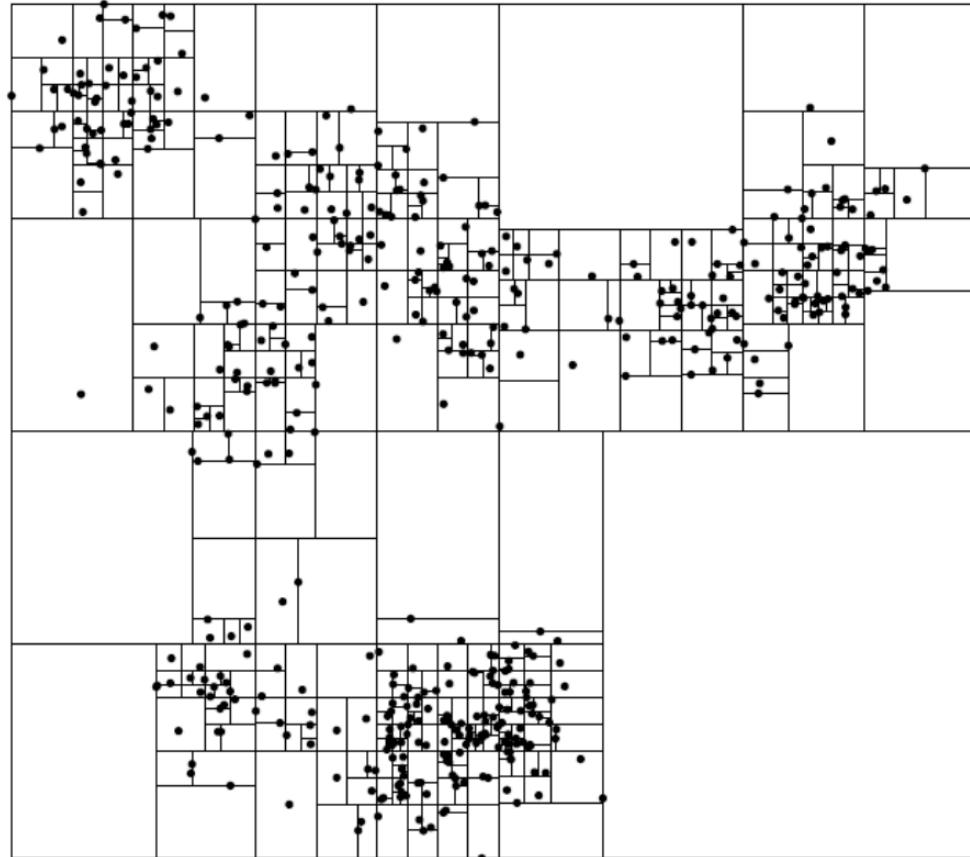
- A randomized algorithm for rapidly finding correspondences between image patches

PatchMatch

- **Problem definition:**
- Given images A and B, for each overlapping patch in image A, compute the offset to the nearest neighbor patch in image B



Previous Work



Time:
 $O(n \log n)$

Kd-tree with PCA

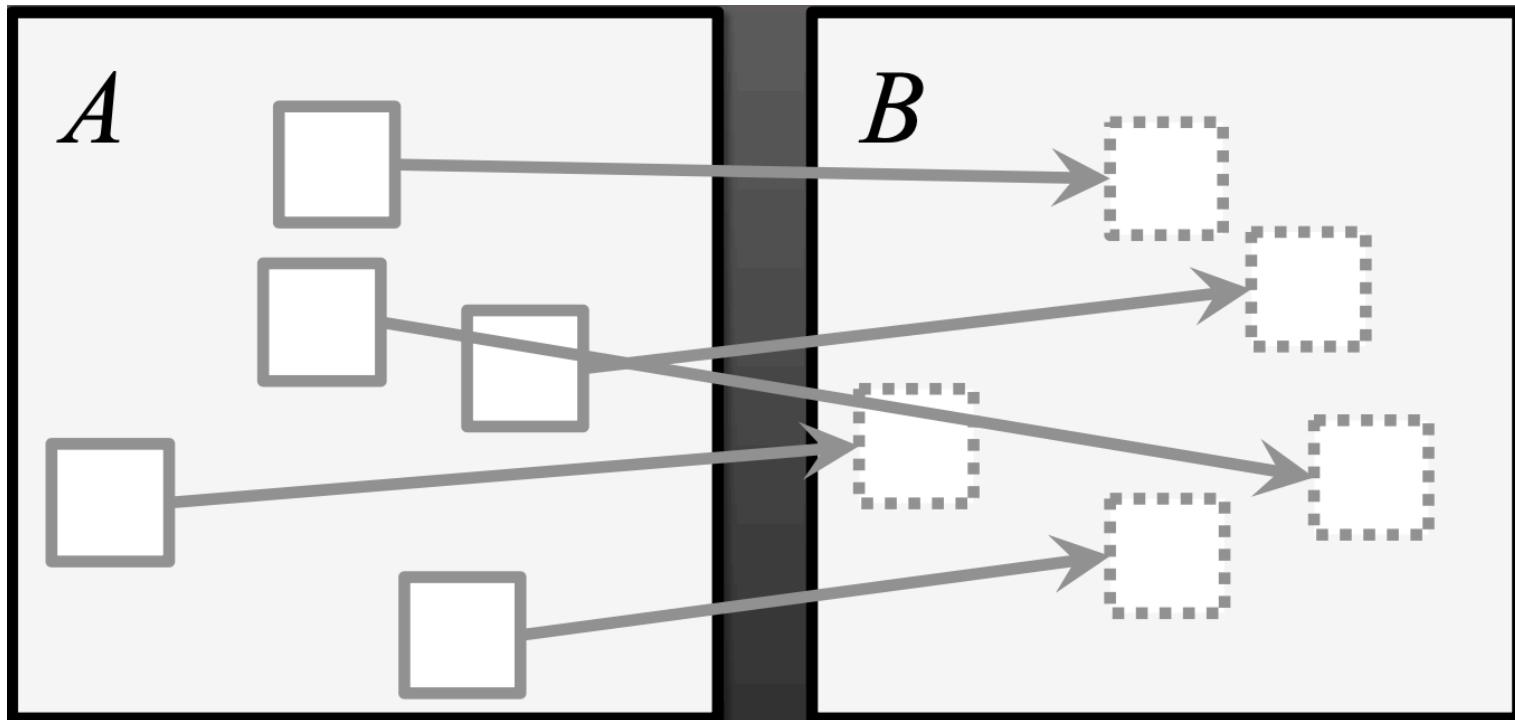
Credit: Hertzmann

Key observation one

- Law of large numbers: a non-trivial fraction of a large field of random offset assignments are likely to be good guesses

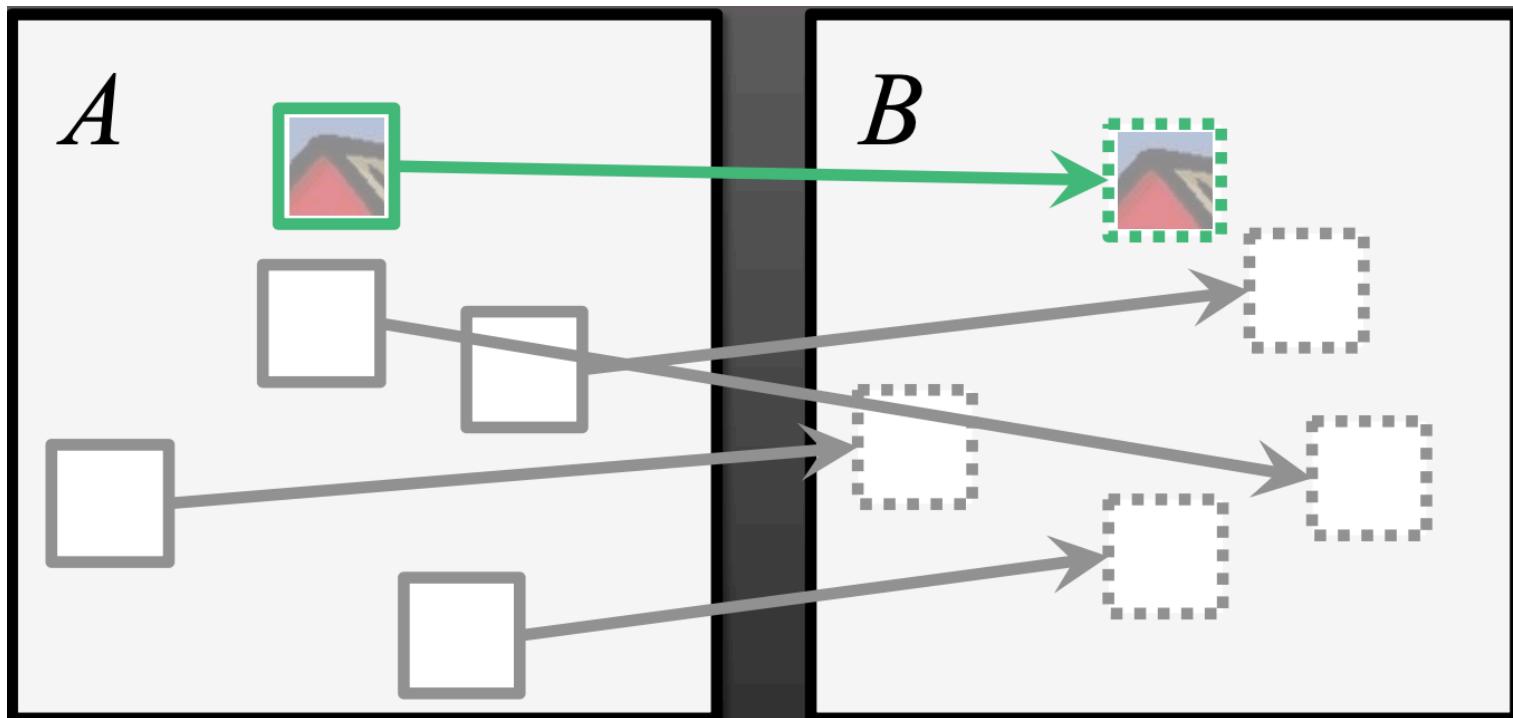
Step 1: Initialization

- Initialization with random values(or derived from prior information)
- $f(x, y) = \text{random value}$



Step 1: Initialization

- Initialization with random values(or derived from prior information)
- $f(x, y) = \text{random value}$



Key observation two: spatial coherence

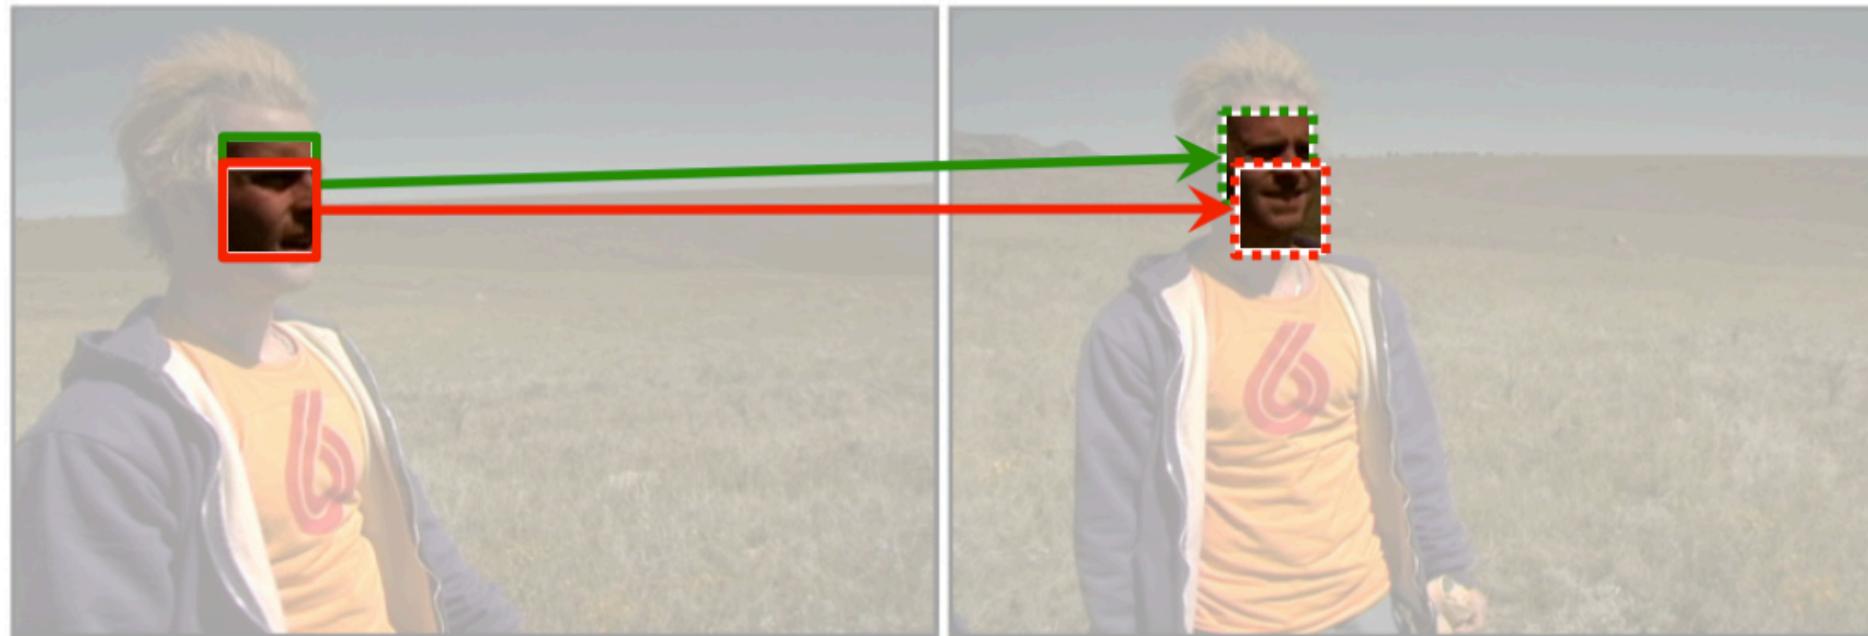
- High coherence of nearest neighbors in natural images
- Nearest neighbor of patch at (x,y) should be a strong hint for where to find nearest neighbor of patch at $(x+1,y)$

Key observation two: spatial coherence



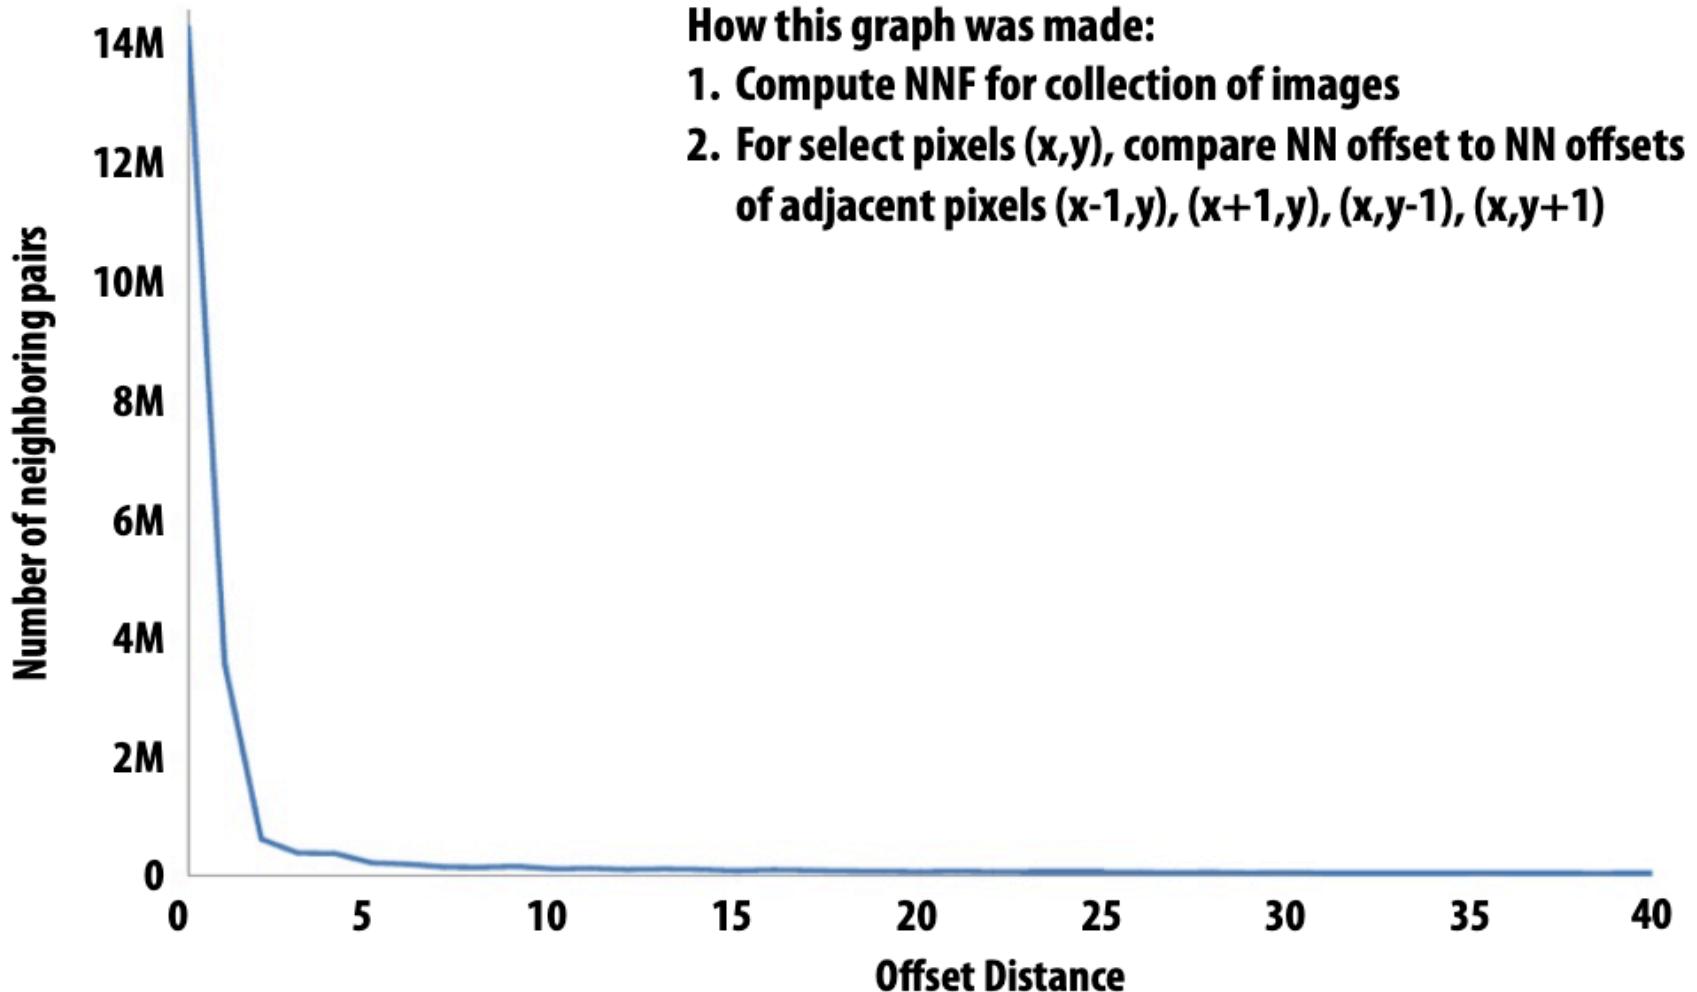
Credit: C. Barnes

Key observation two: spatial coherence



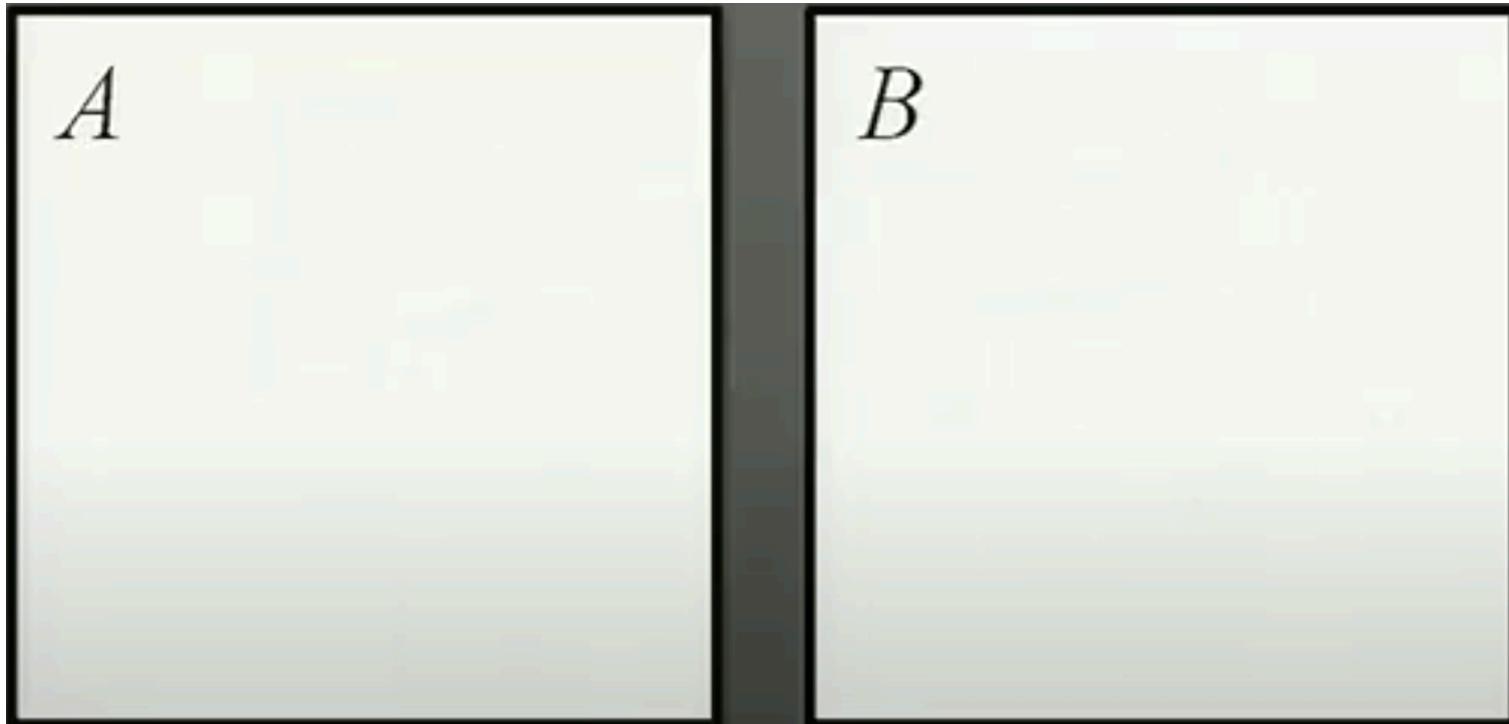
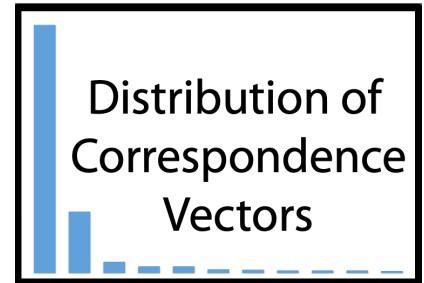
Credit: C. Barnes

Use Statistics



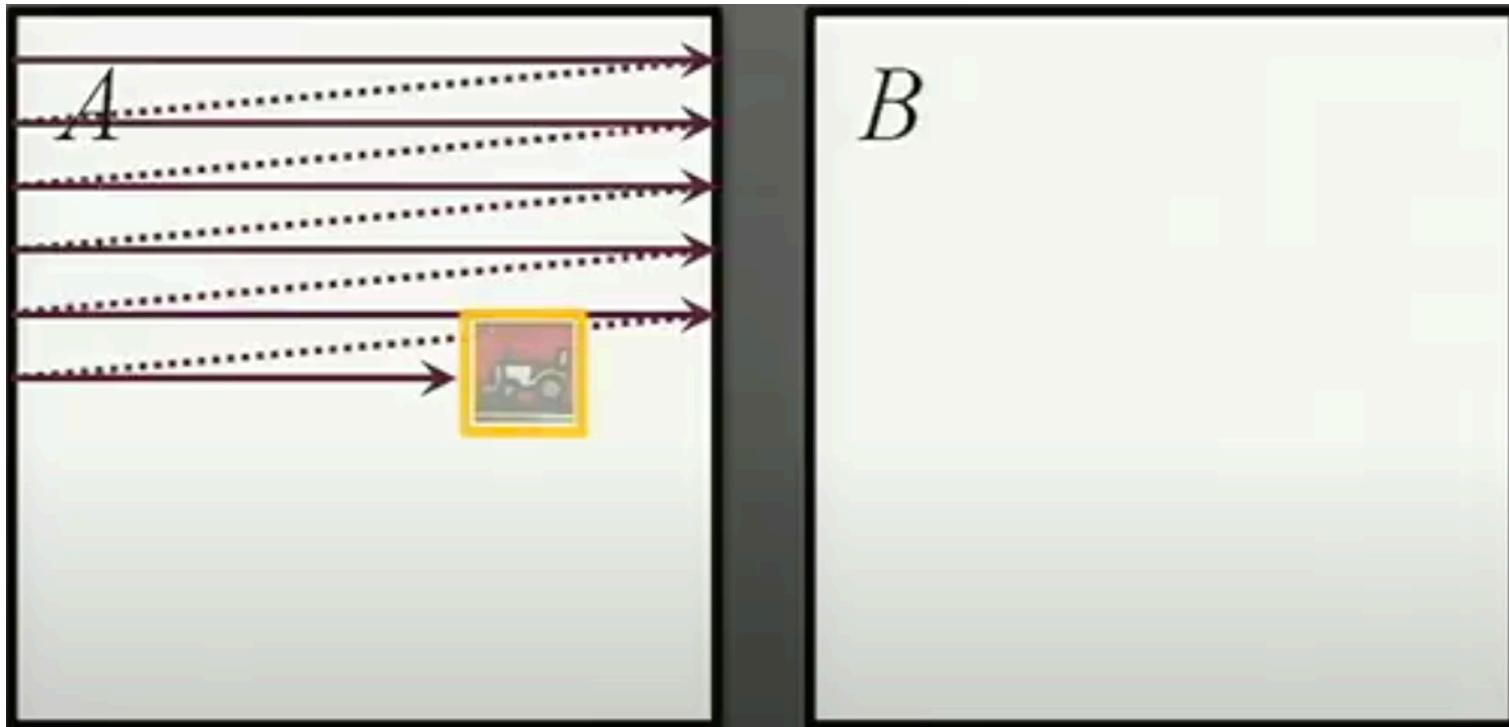
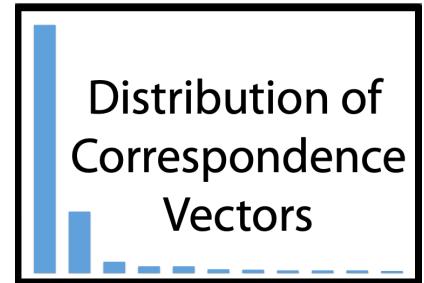
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor(or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



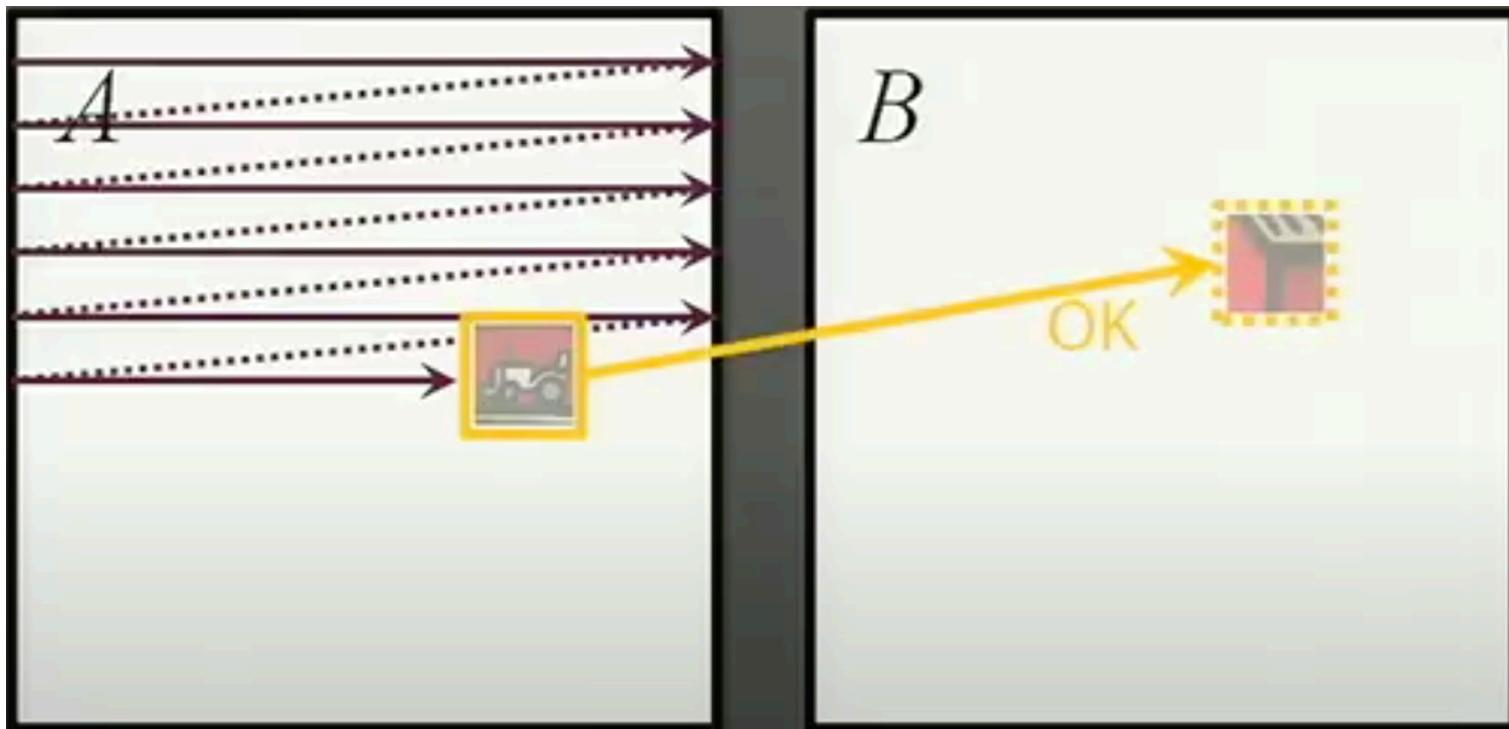
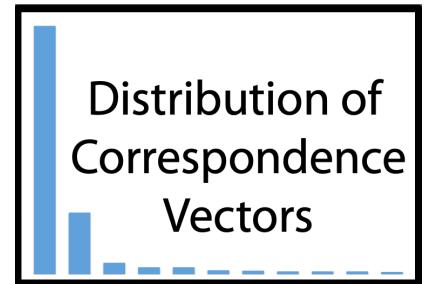
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor(or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



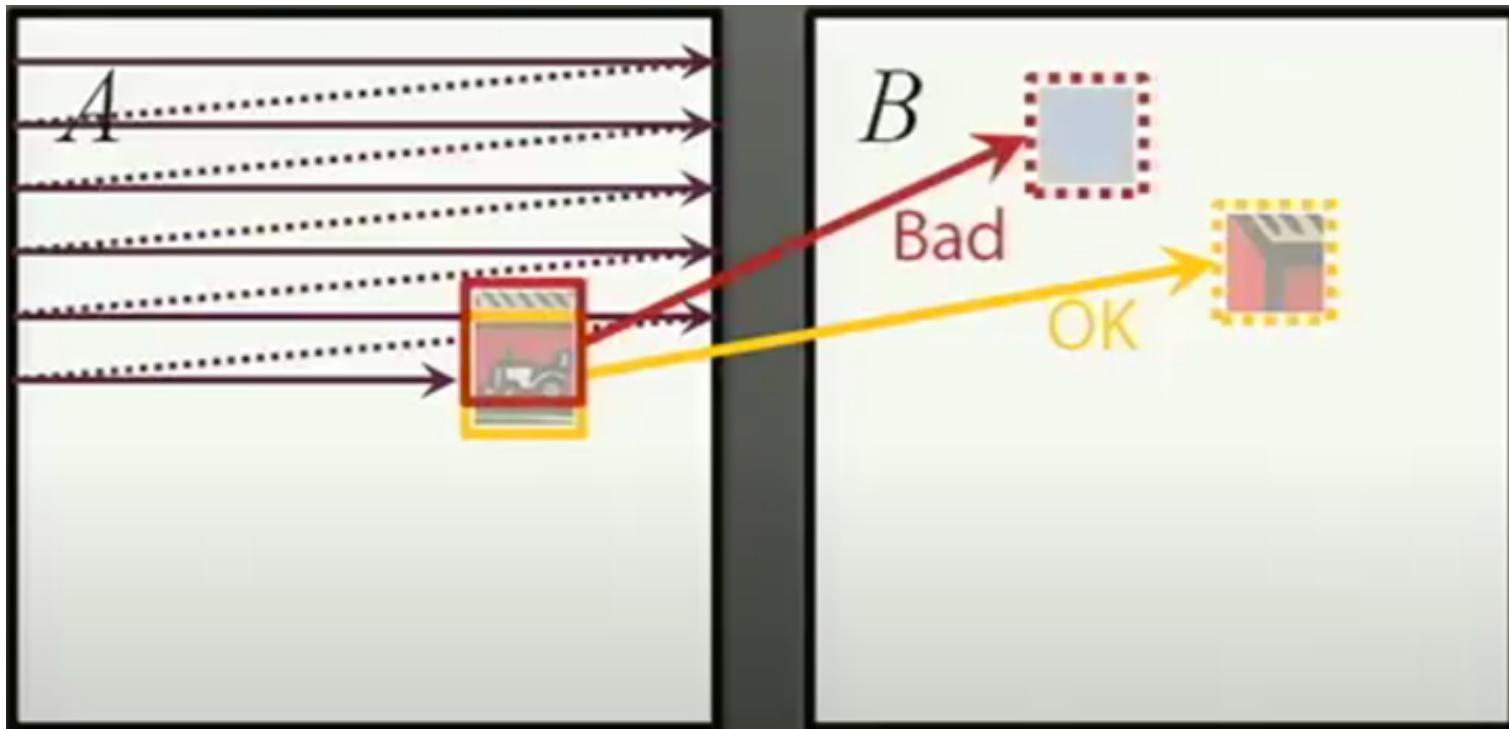
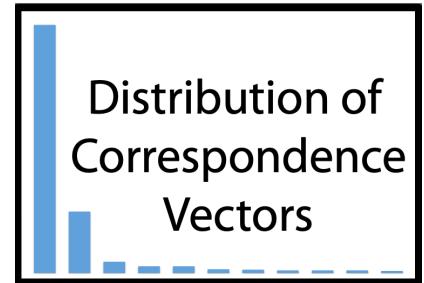
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor (or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



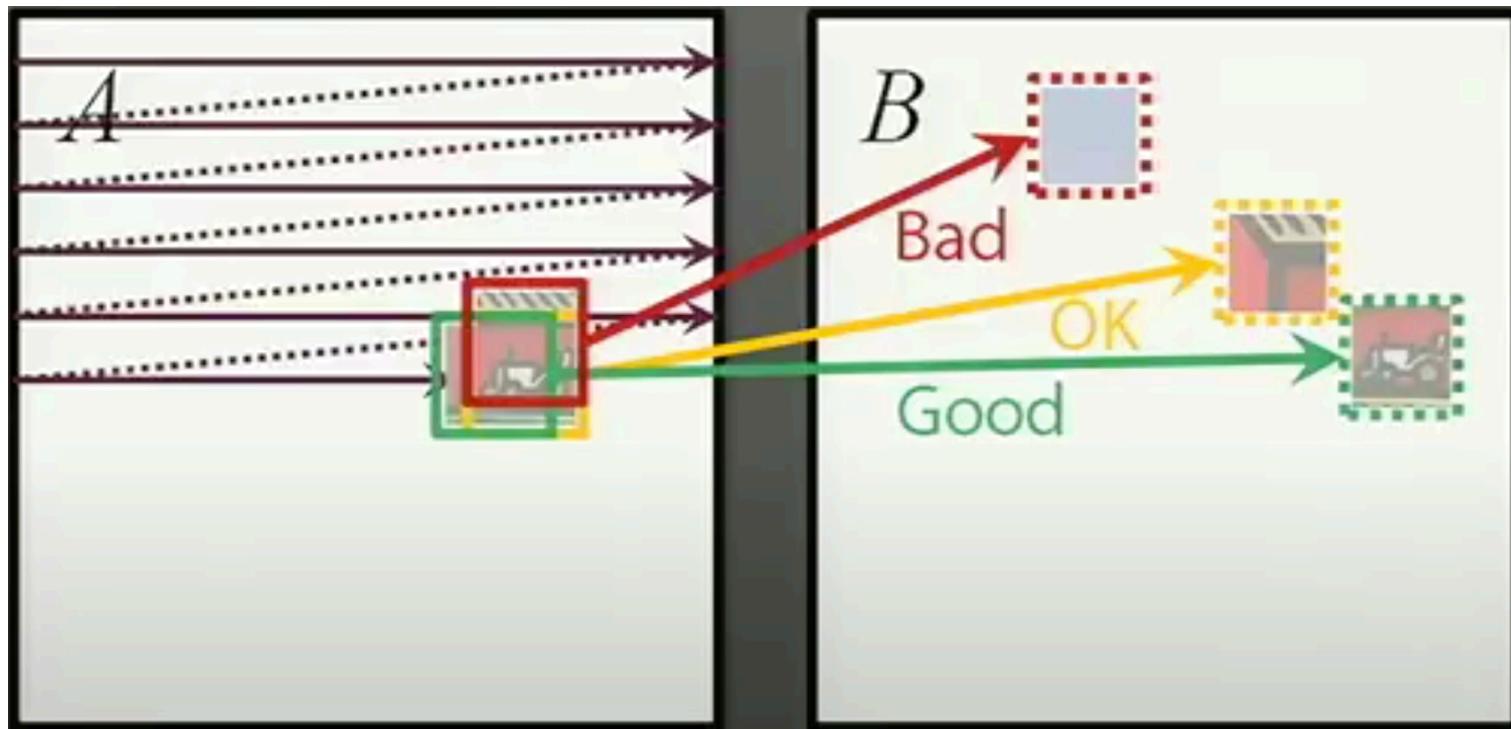
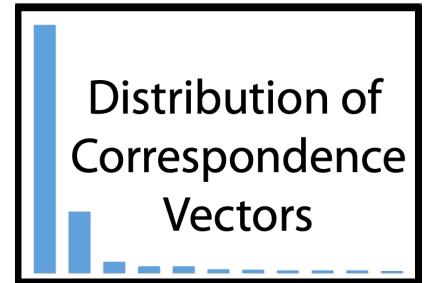
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor (or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



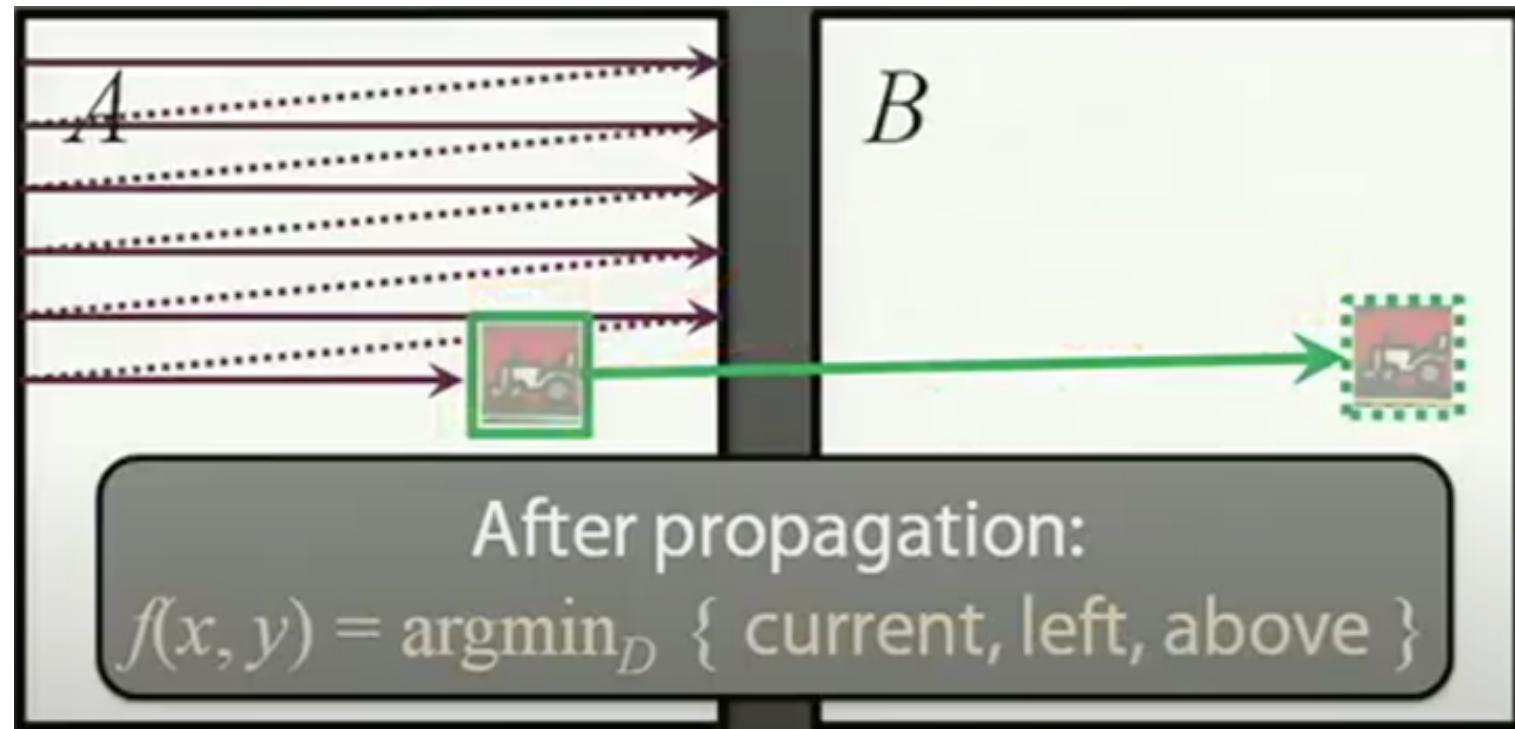
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor (or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



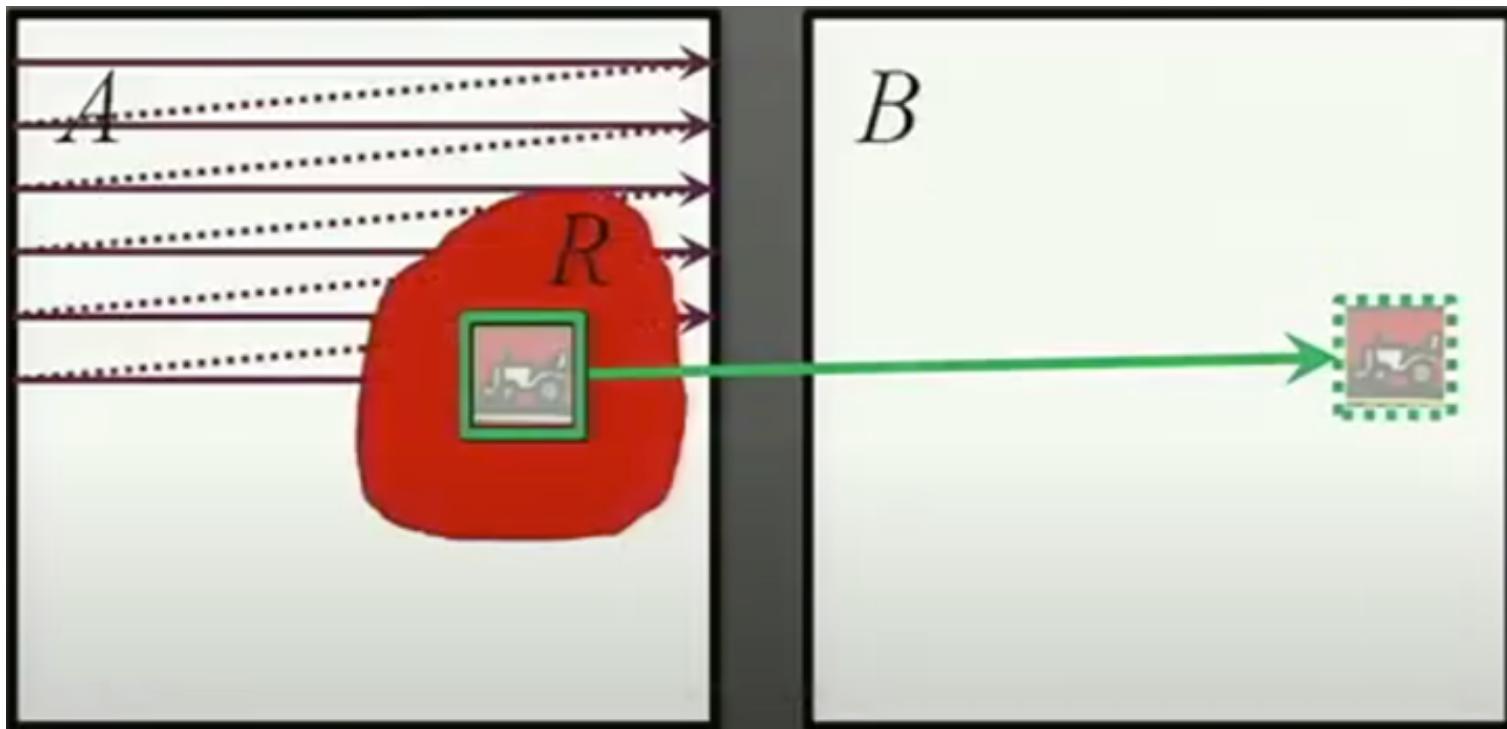
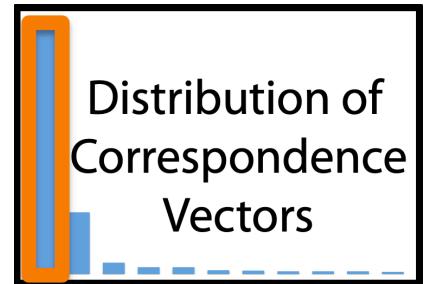
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor (or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



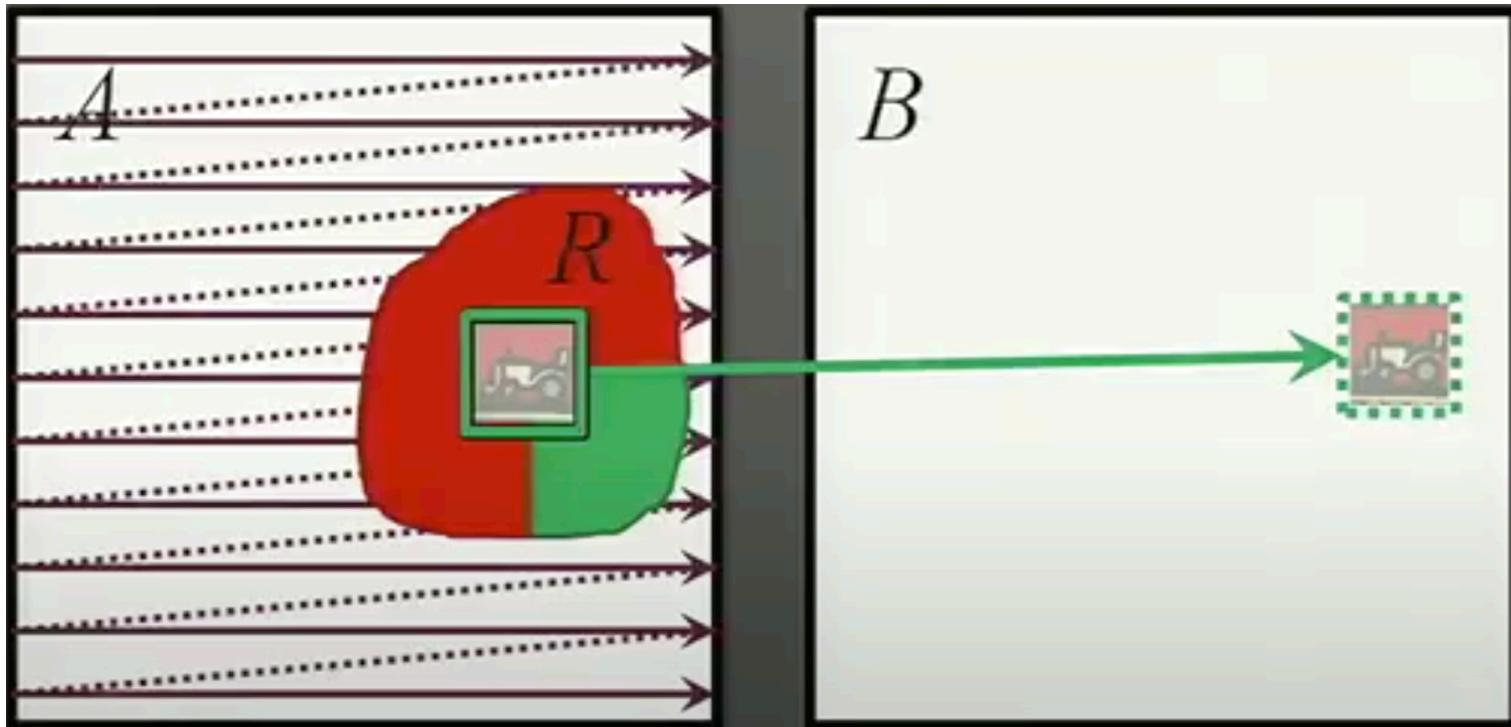
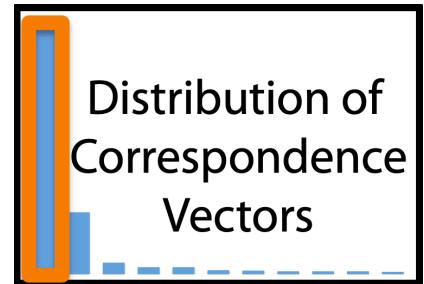
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor (or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



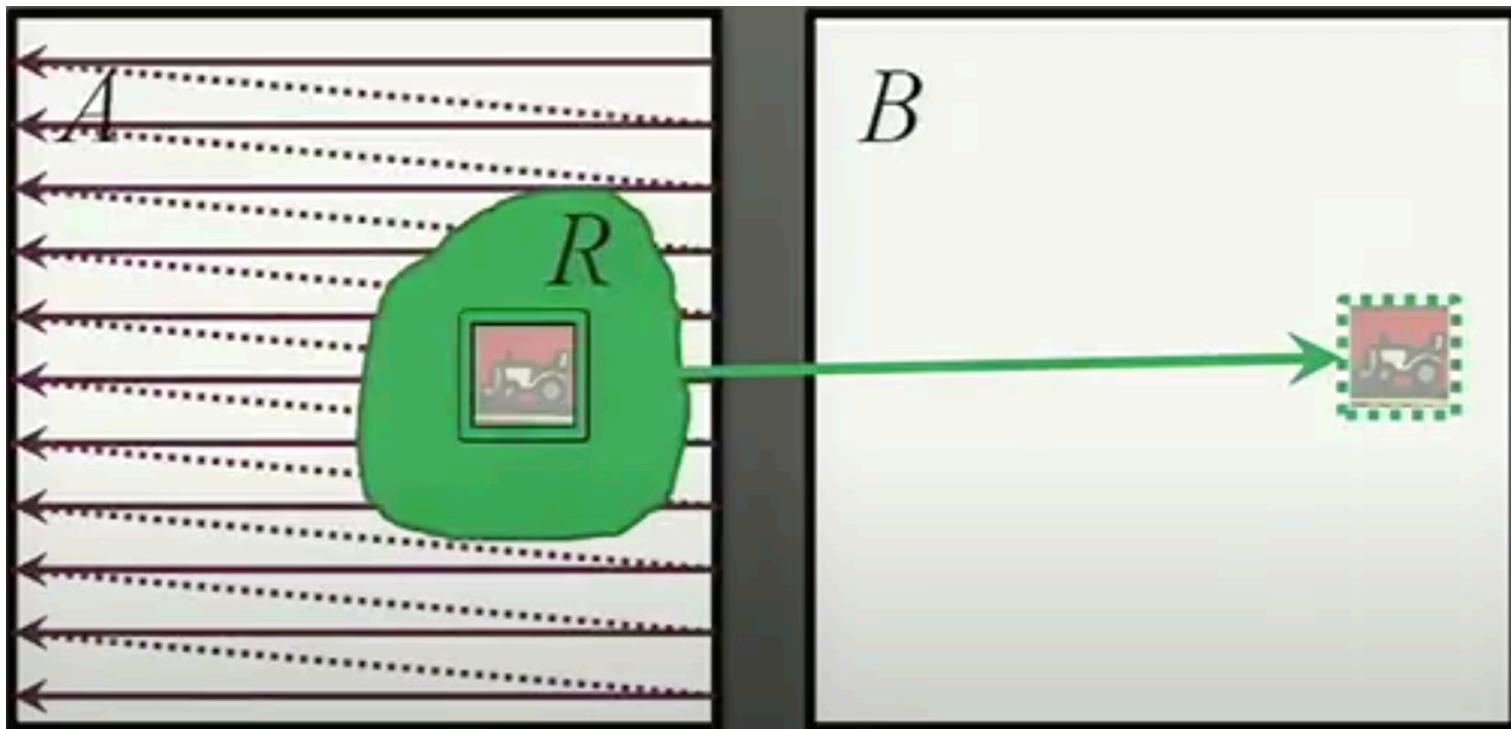
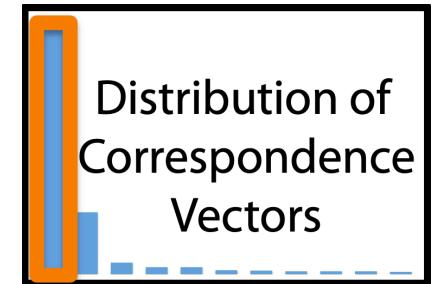
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor (or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x - 1, y), f(x, y - 1))$



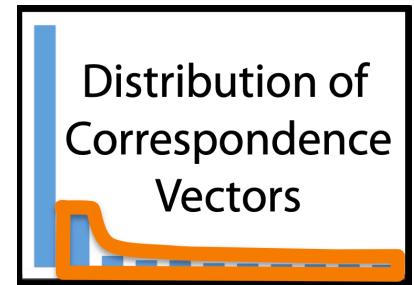
Step 2: Propagation

- Try to improve offset estimate by exploiting spatial coherence with left and top neighbor (or right, bottom)
- $f(x, y) = \operatorname{argmin}_d(f(x, y), f(x + 1, y), f(x, y + 1))$



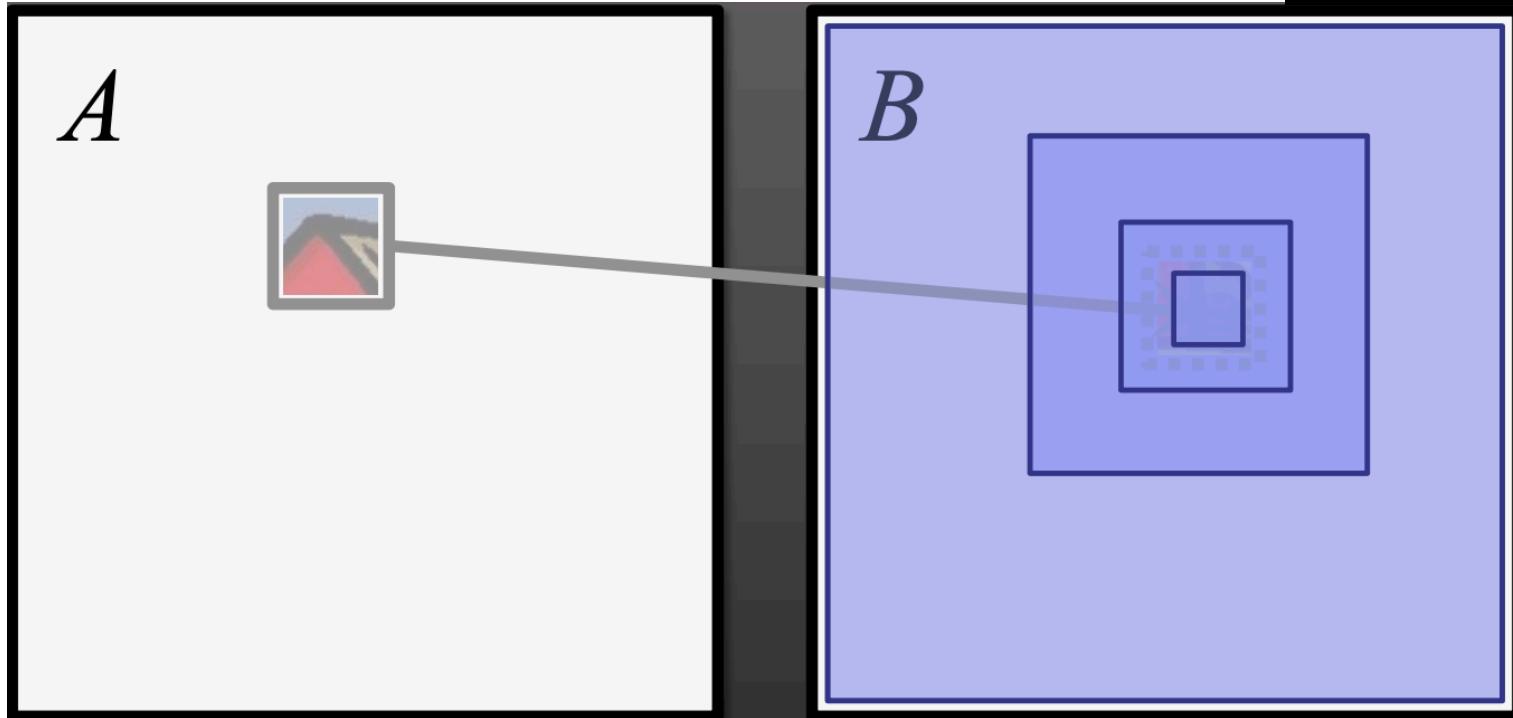
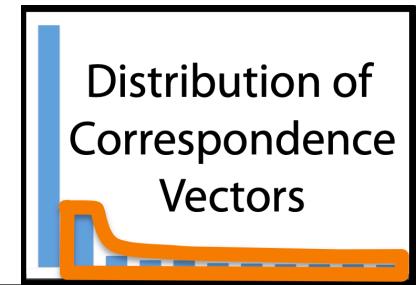
Step 3: Random Search

- **Avoiding local minima**
- **Random search in the neighborhood of the best offset found so far.**
- $f(x, y) = \operatorname{argmin}_d \{\text{candidate correspondence}\}$



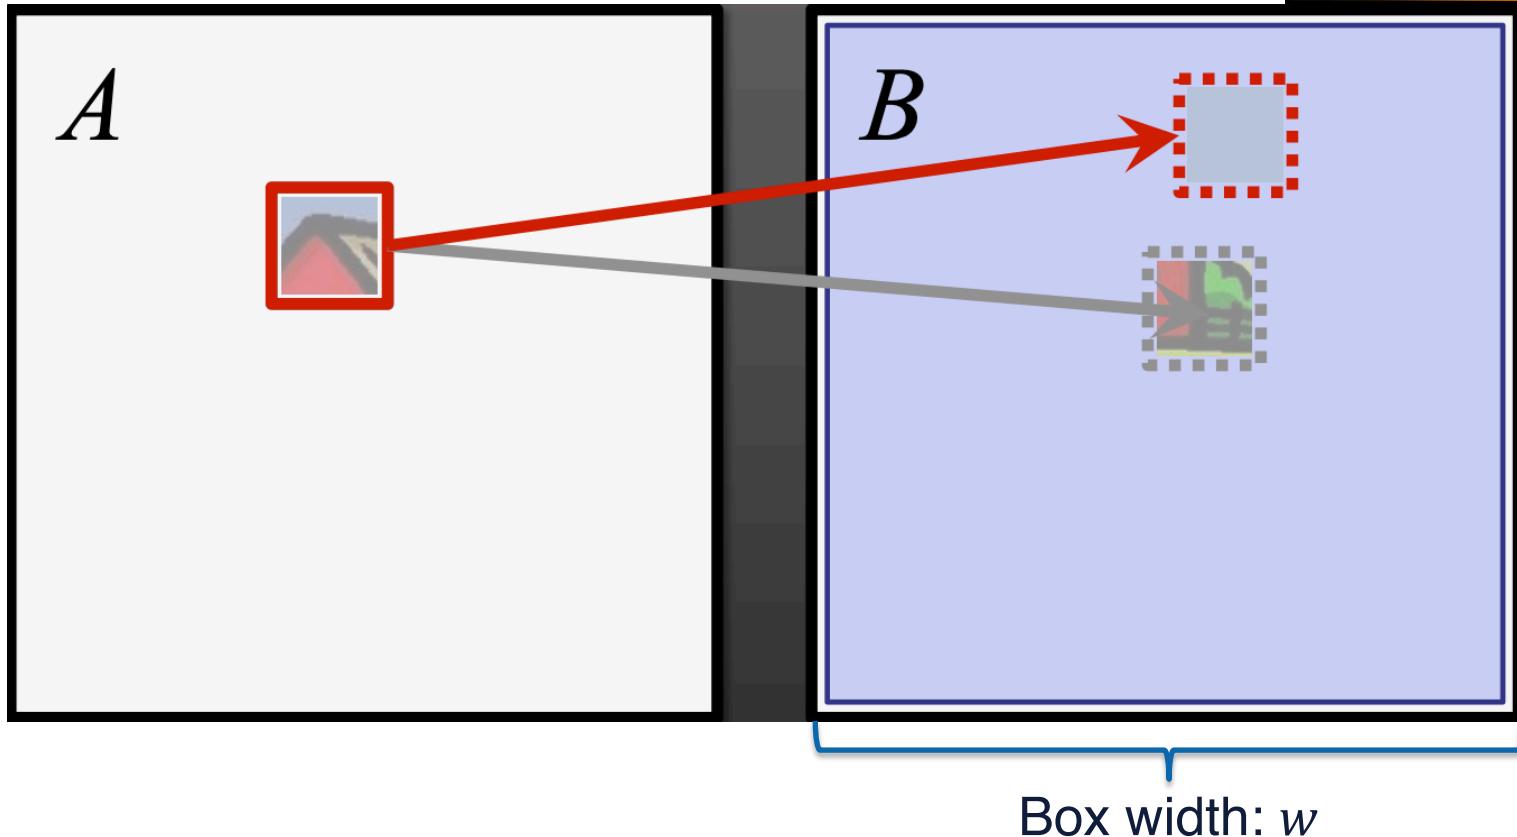
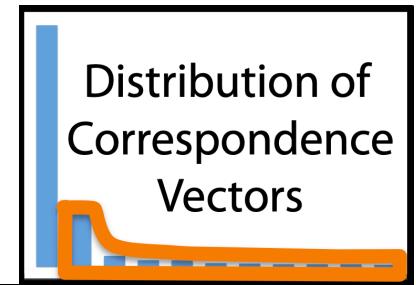
Step 3: Random Search

- random search in the neighborhood of the best offset found so far.
- $f(x, y) = \operatorname{argmin}_d \{\text{candidate correspondence}\}$



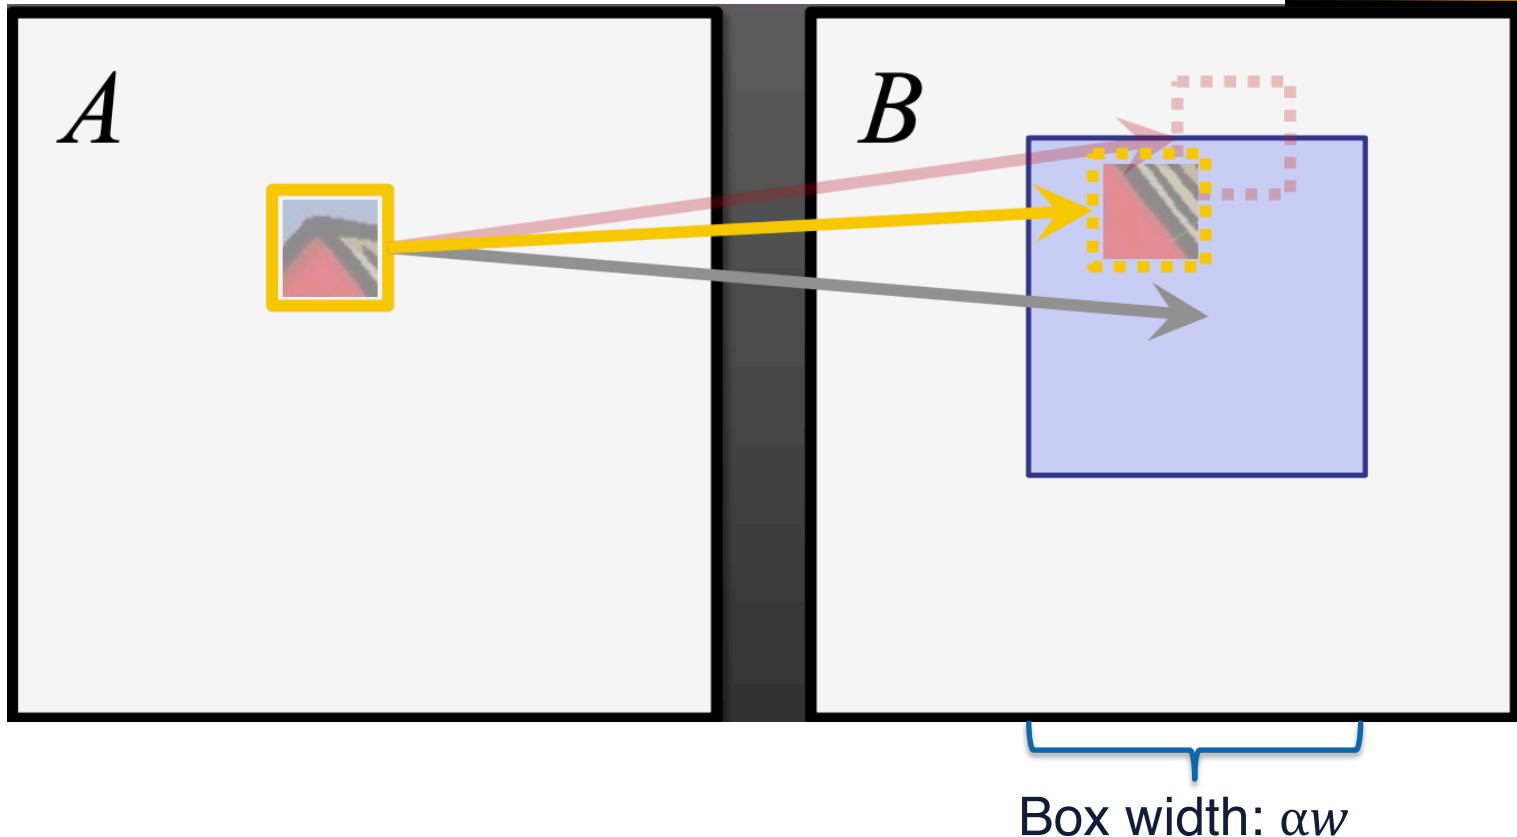
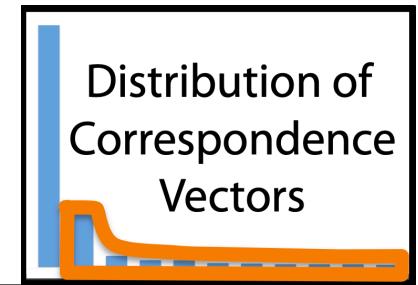
Step 3: Random Search

- random search in the neighborhood of the best offset found so far.
- $f(x, y) = \operatorname{argmin}_d \{ \text{candidate correspondence} \}$



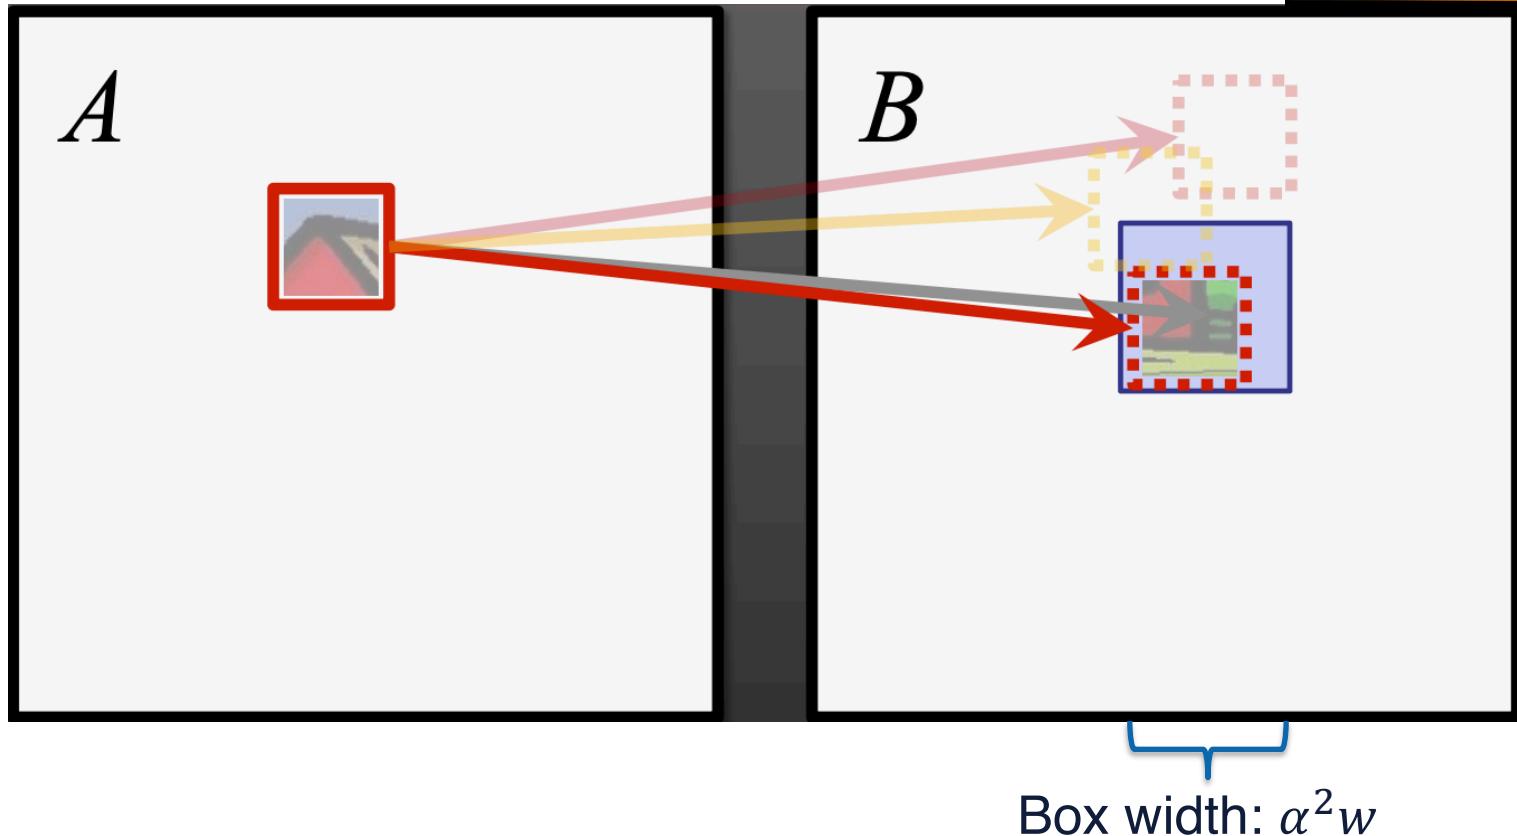
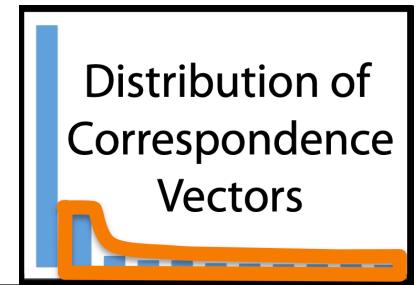
Step 3: Random Search

- random search in the neighborhood of the best offset found so far.
- $f(x, y) = \operatorname{argmin}_d \{\text{candidate correspondence}\}$



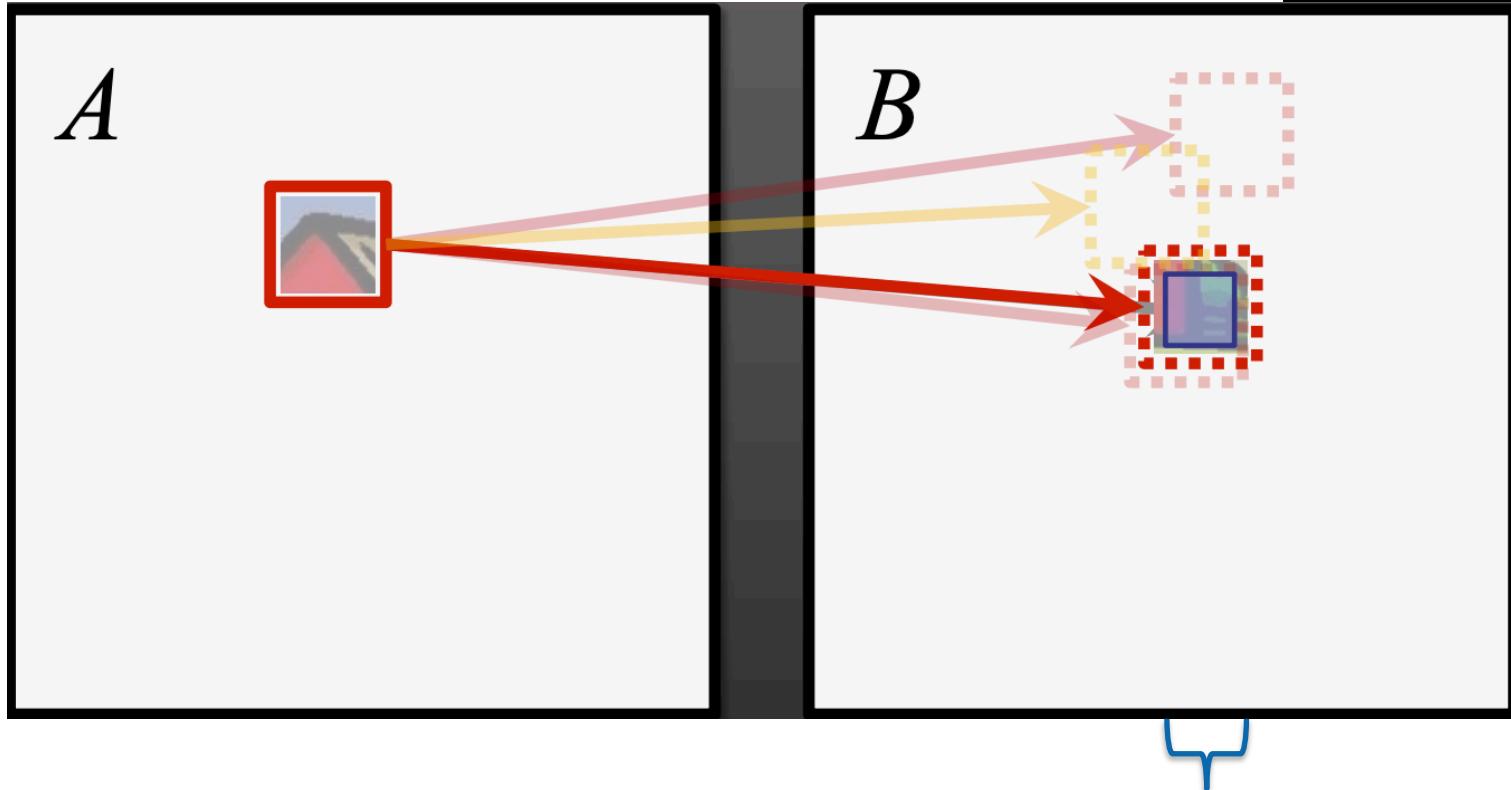
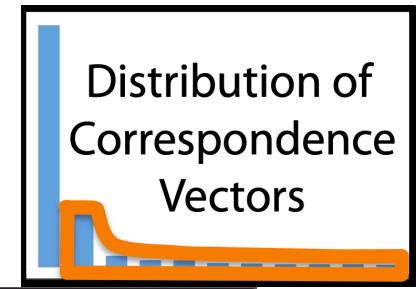
Step 3: Random Search

- random search in the neighborhood of the best offset found so far.
- $f(x, y) = \operatorname{argmin}_d \{\text{candidate correspondence}\}$



Step 3: Random Search

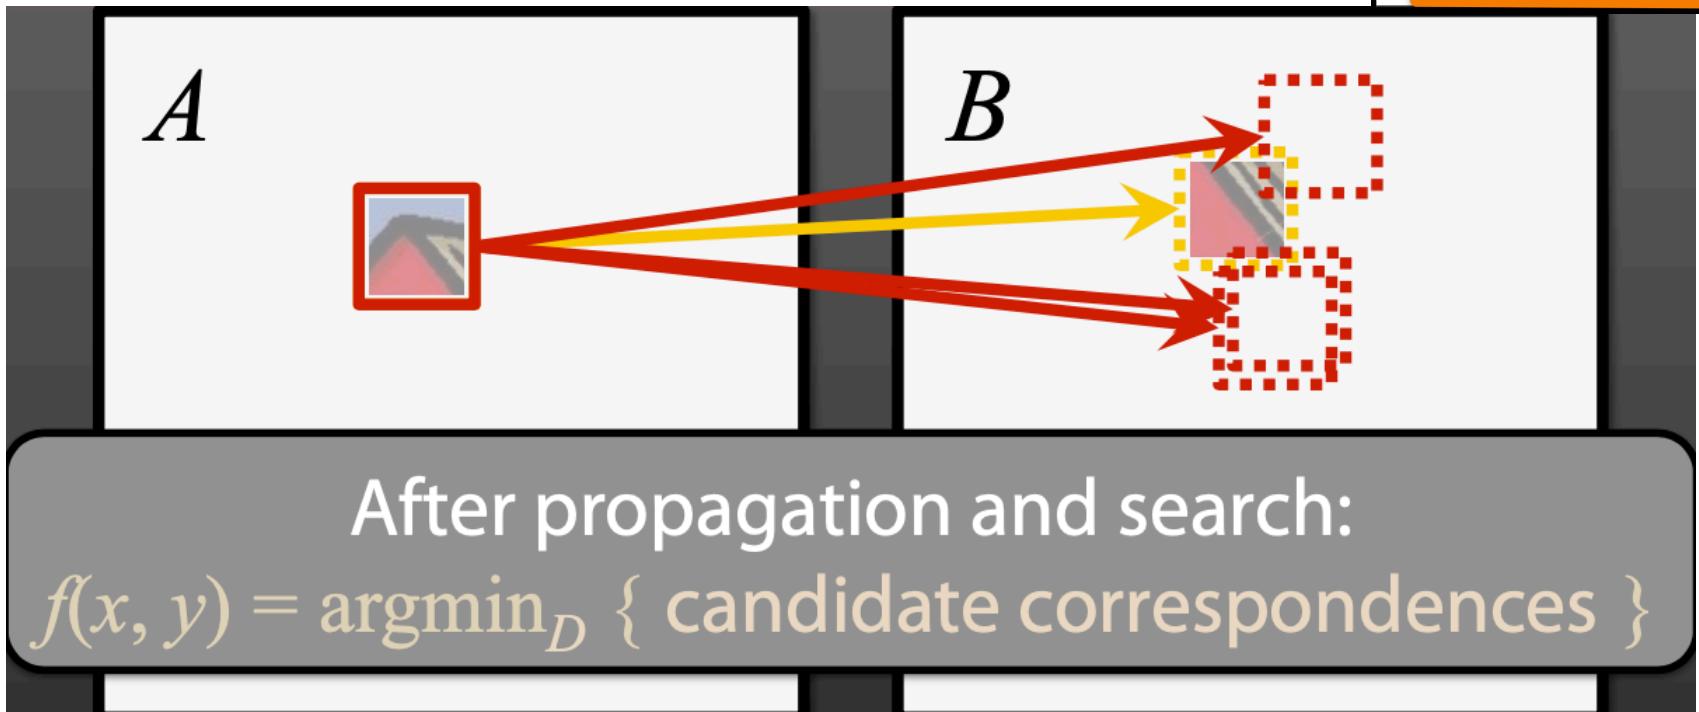
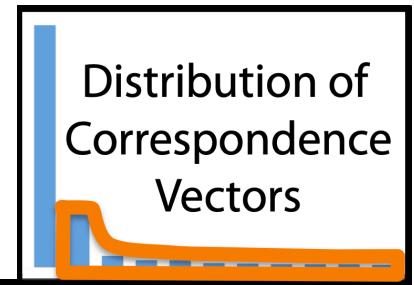
- random search in the neighborhood of the best offset found so far.
- $f(x, y) = \operatorname{argmin}_d \{\text{candidate correspondence}\}$



Box width: 1 pixel

Step 3: Random Search

- random search in the neighborhood of the best offset found so far.
- $f(x, y) = \operatorname{argmin}_d \{ \text{candidate correspondence} \}$



Summary

- PatchMatch:
 - Step 1: Initialization
 - Step 2: Propagation
 - Step 3: Random Search



Repeat
until
converged

Summary

- PatchMatch:
 - Step 1: Initialization
 - Step 2: Propagation
 - Step 3: Random Search
 - key insights:
 - some good patch matches can be found via random sampling.
 - natural coherence in the imagery allows us to propagate such matches quickly to surrounding areas.
- 
- Repeat
until
converged

Experiment:
Reconstruct A using
patches from B

Image A

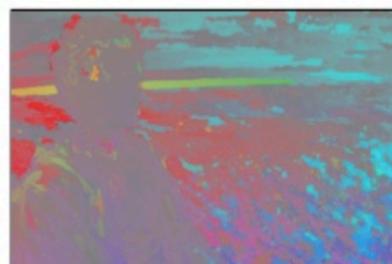


**Image B
(source of
patches)**



Random init:

1/4 through iter 1



End of iter 1

Iter 2

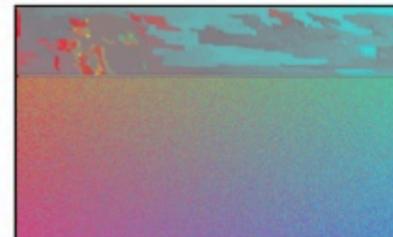
Iter 5

Credit: Barnes

Image A



**Image B
(source of
patches)**



**Experiment:
Reconstruct A using
patches from B**

10-100x faster than kd-tree!



End of iter 1



Iter 2



Iter 5

Credit: Barnes

Why does it work?

- Assume source and target images have equal size (M pixels) and that random initialization is used.
- The odds of any one location being assigned the best offset: $1 / M$
- But for M pixels:
 - The odds of at least one offset being correctly assigned are quite good:

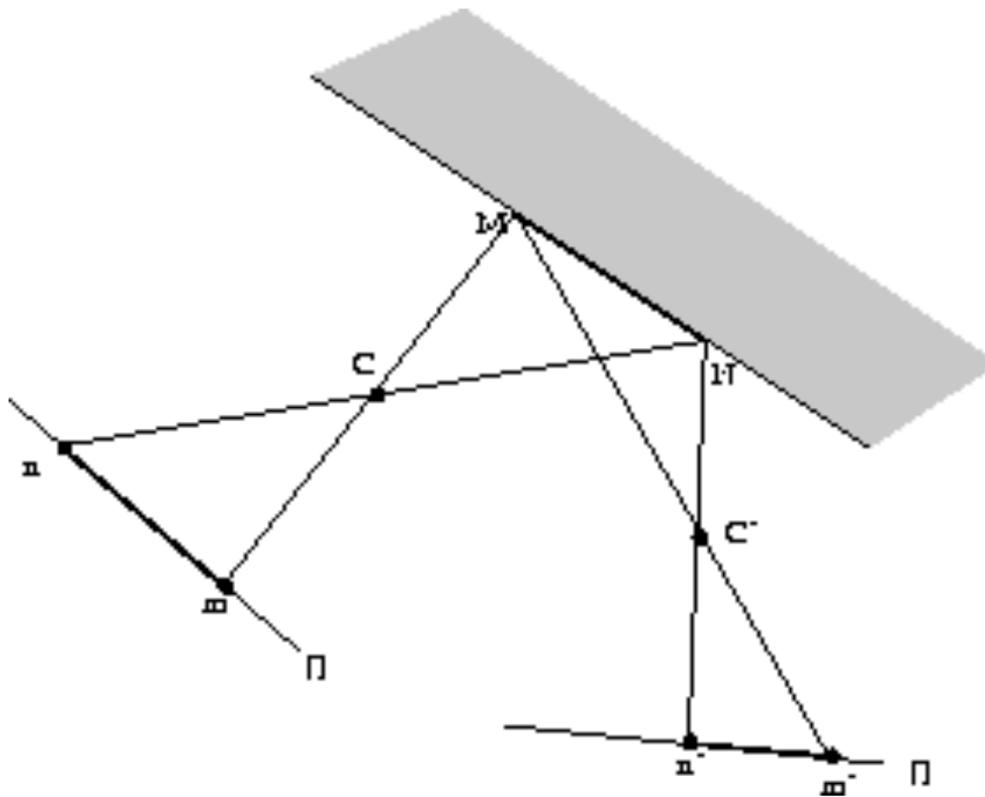
$$1 - \left(1 - \frac{1}{M}\right)^M \quad \text{E.g. } M=10\text{e}5, \text{ this is } (1 - 0.367)$$

If top C nearest neighbors are enough, the odds will be $1 - \left(1 - \frac{C}{M}\right)^M$

PatchMatch Stereo

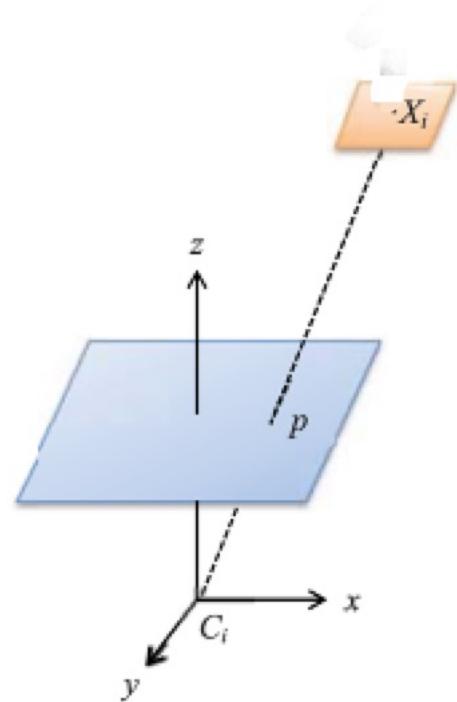
PatchMatch Stereo

Figure 6: Foreshortening due to the change of viewing position and direction.



PatchMatch Stereo

- Extend to find an approximate nearest neighbor according to a **plane**.
- Offset -> depth

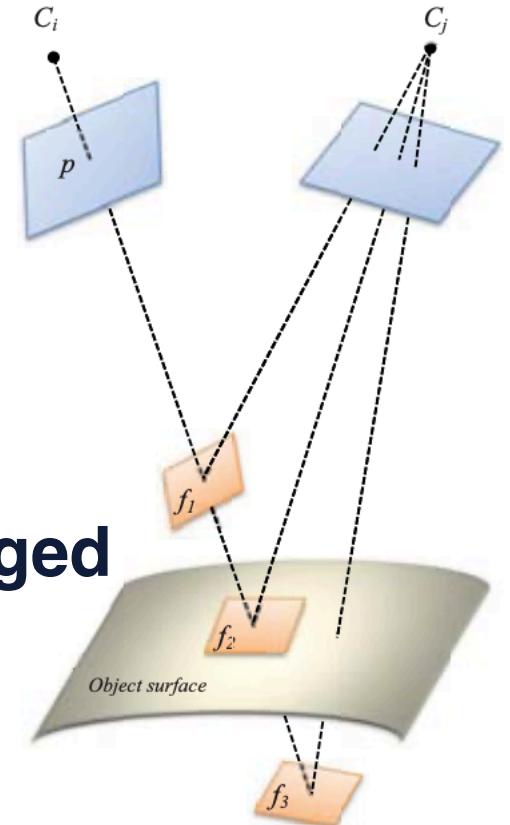


PatchMatch Stereo

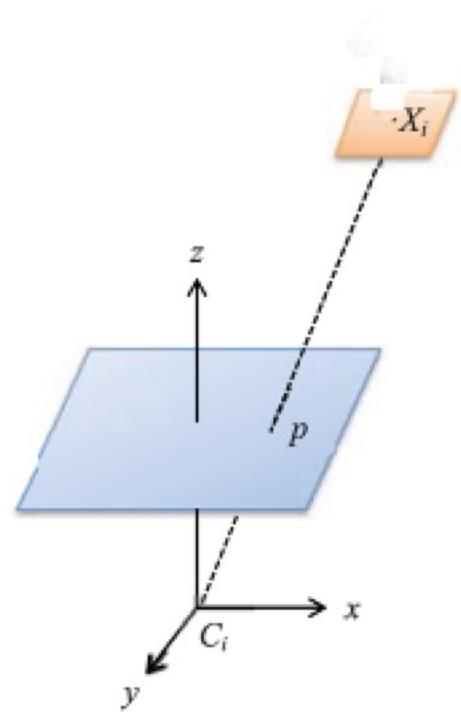
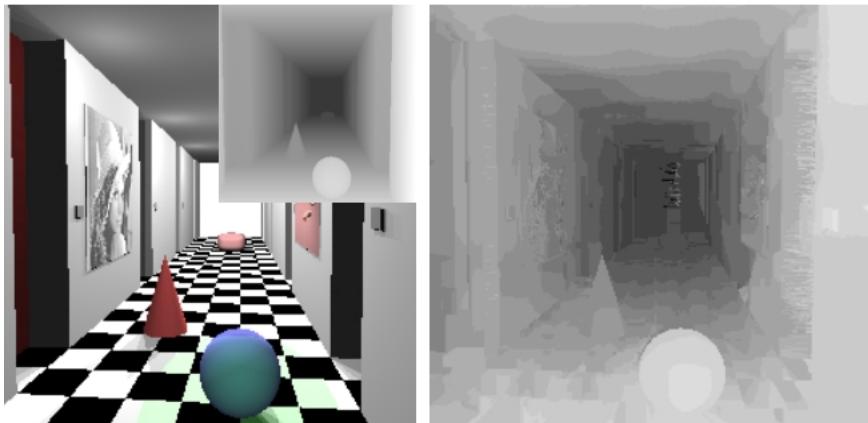
- Step 1: Initialization
- Step 2: Propagation
- Step 3: Random Search



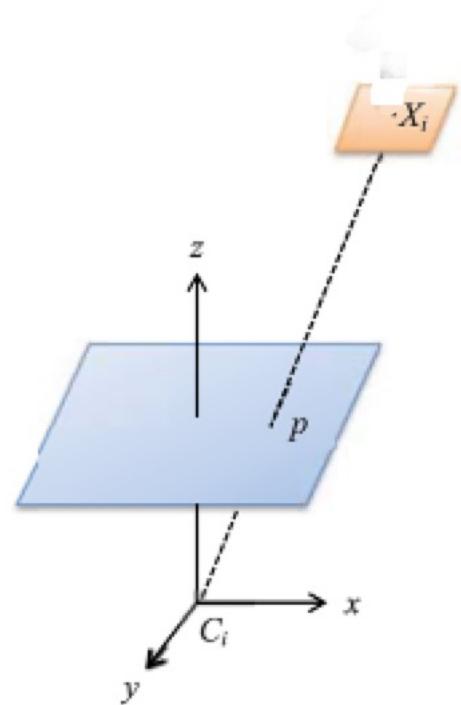
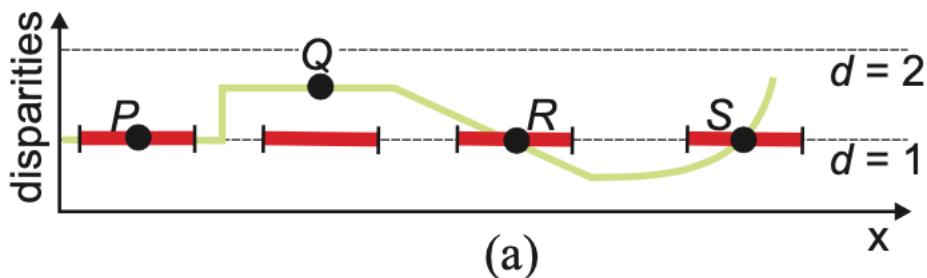
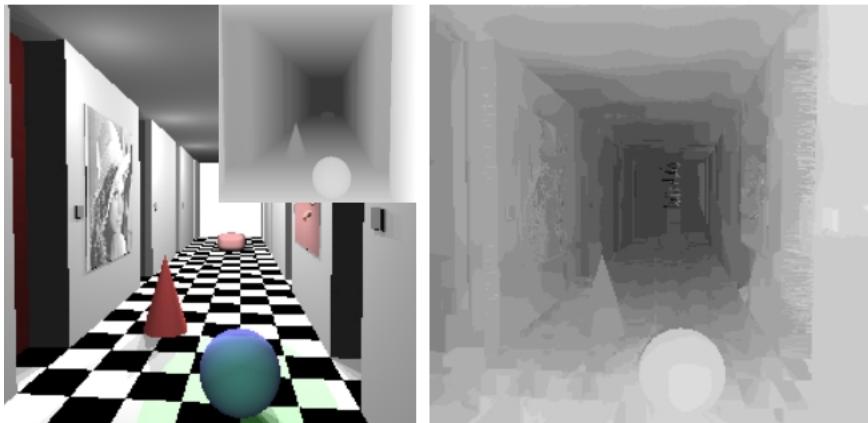
Repeat
until
converged



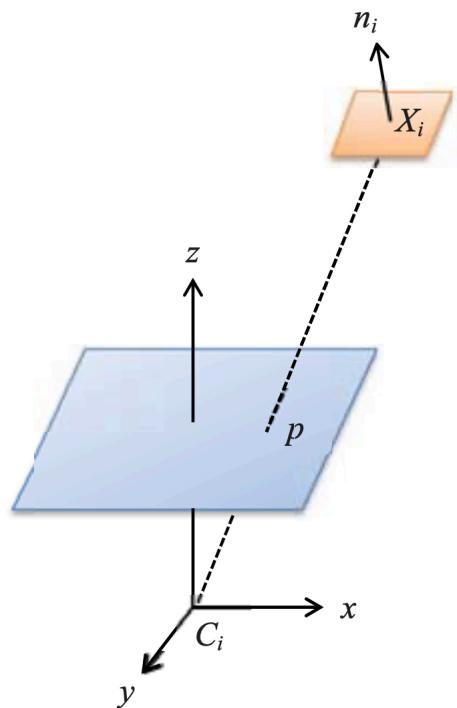
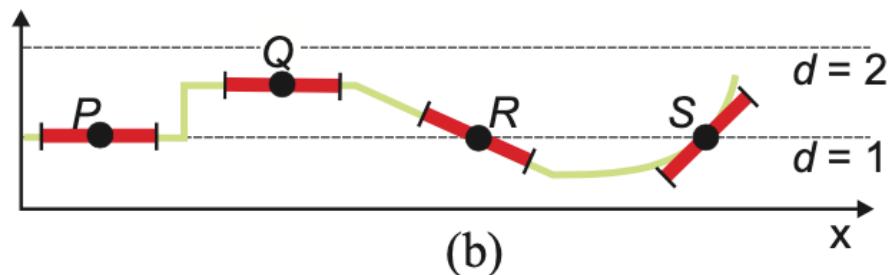
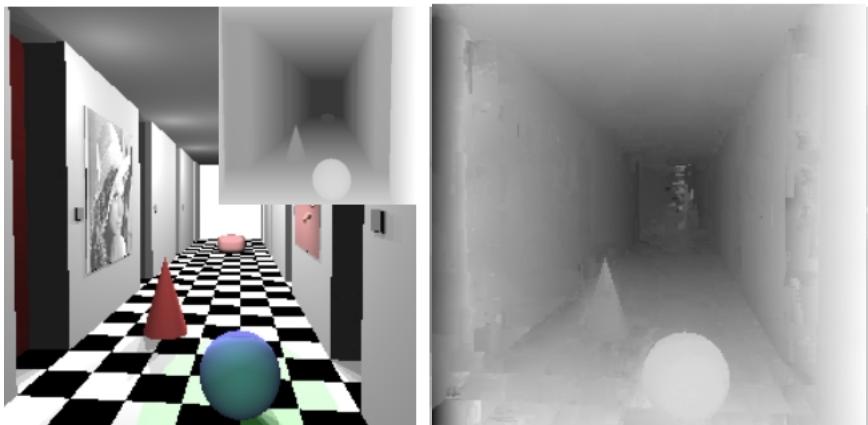
PatchMatch Stereo



PatchMatch Stereo

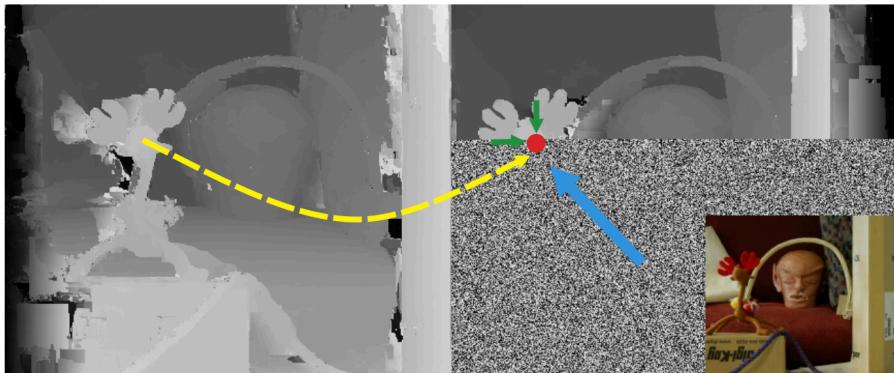


PatchMatch Stereo



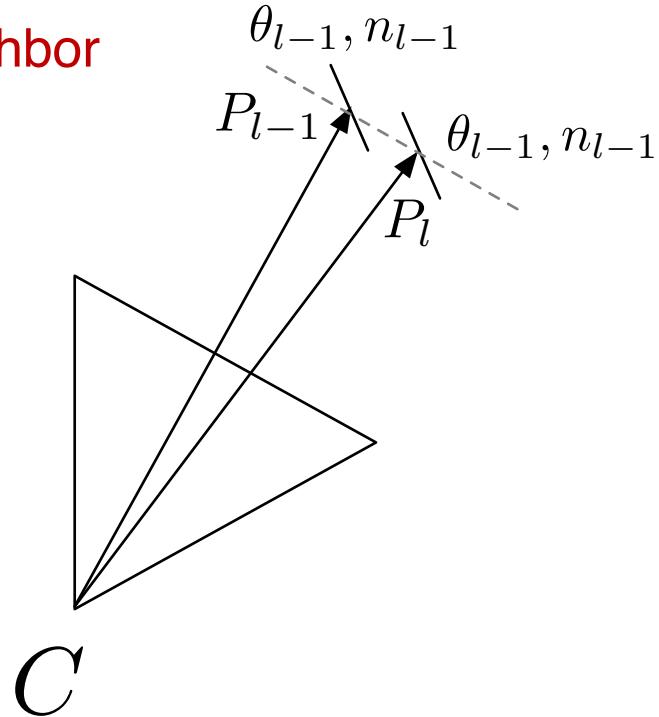
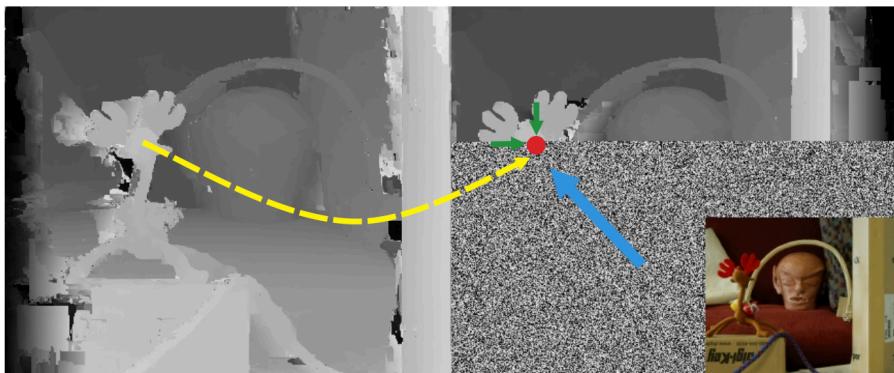
PatchMatch Stereo

- For Each Pixel
 - Assign Random Depth and Normal
- For N Iterations
 - For Each Pixel
 - Propagate Depth and Normal From Neighbor
 - Sample New Random Depth and Normal
 - Update Depth



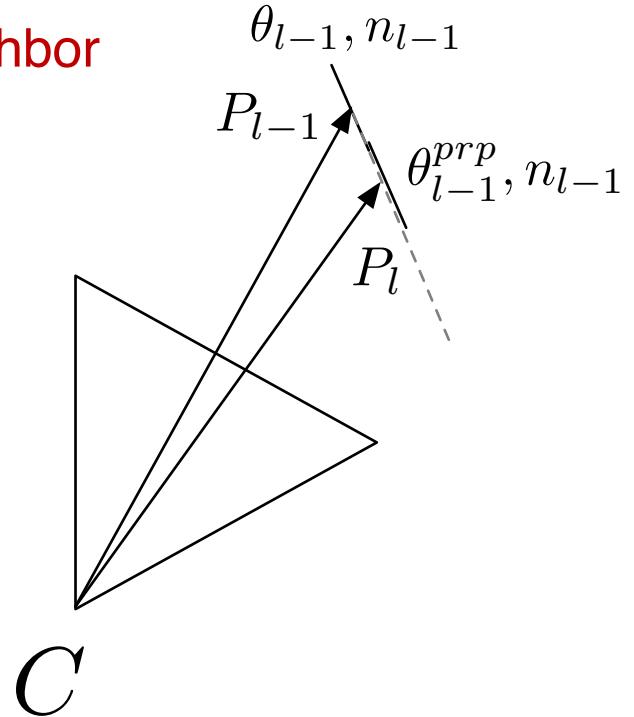
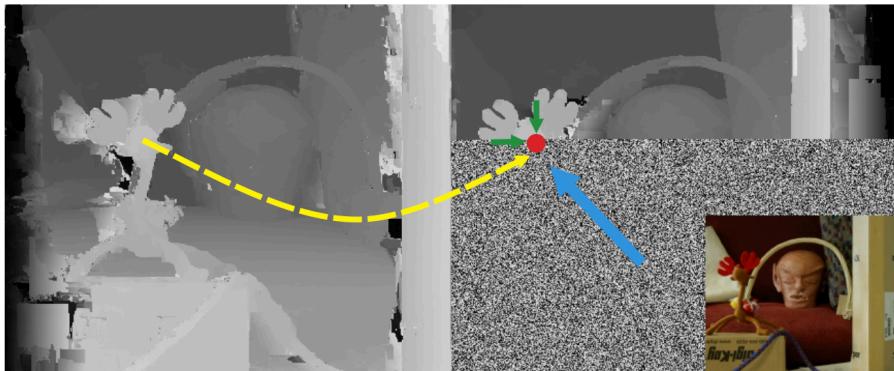
PatchMatch Stereo

- For Each Pixel
 - Assign Random Depth and Normal
- For N Iterations
 - For Each Pixel
 - Propagate Depth and Normal From Neighbor
 - Sample New Random Depth and Normal
 - Update Depth



PatchMatch Stereo

- For Each Pixel
 - Assign Random Depth and Normal
- For N Iterations
 - For Each Pixel
 - Propagate Depth and Normal From Neighbor
 - Sample New Random Depth and Normal
 - Update Depth



Summary

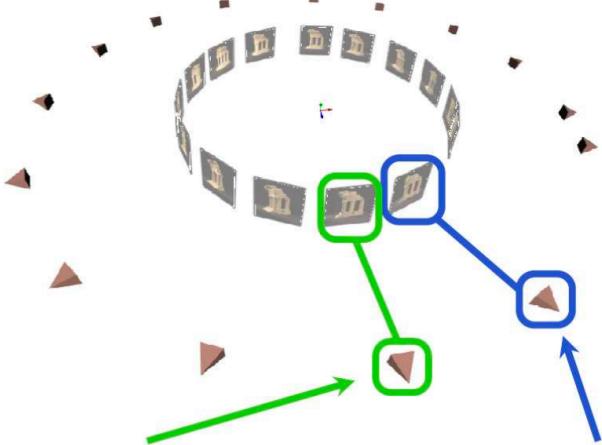
- **Problem Definition:**
 - Finding a “good” slanted support plane at each pixel.
- The difference with vanilla PatchMatch
 - (offset) -> (depth, normal)

View Selection

- Step 1: Source view selection
- Step 2: Depth-map computation
- Step 3: Depth-map merging



Key steps:
1. How to chose source images
2. How to compute depth map

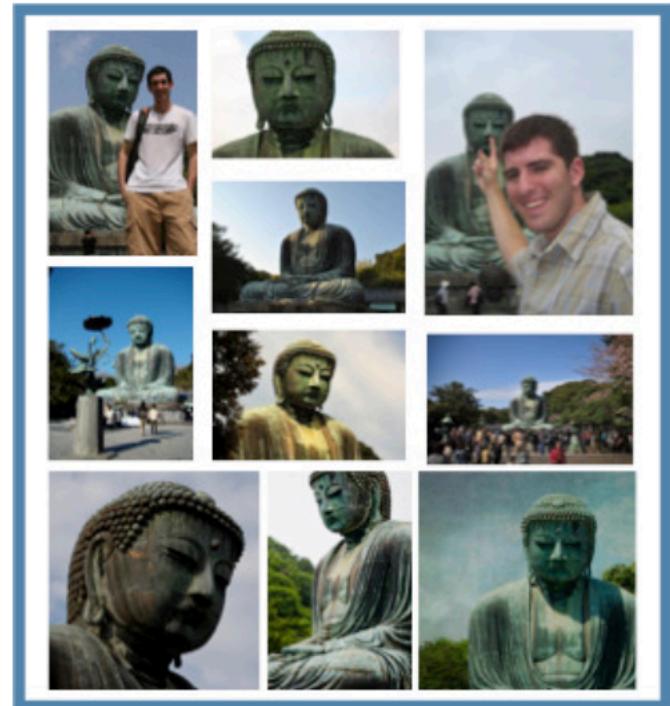


View Selection

- *How to robustly integrate photo-consistency measurements from multiple views?*



Reference image

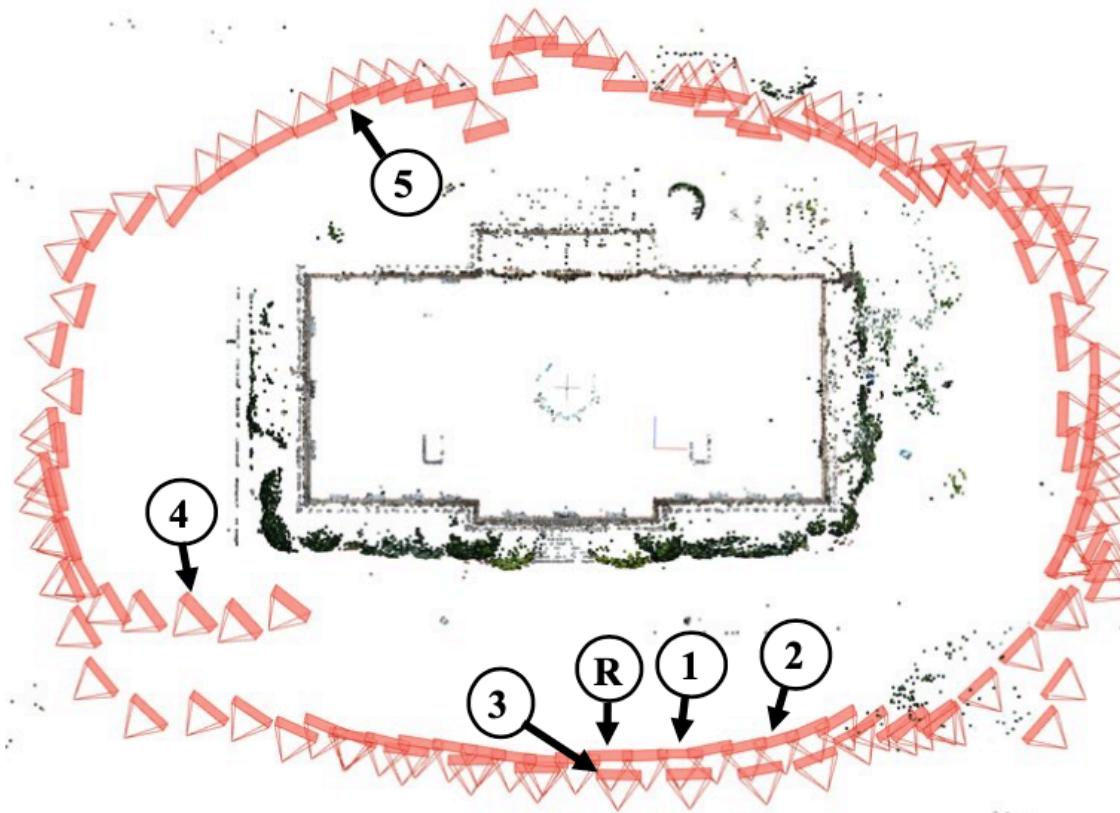


211 Source images (only 10 are shown)

Credit: E. Dunn

View Selection

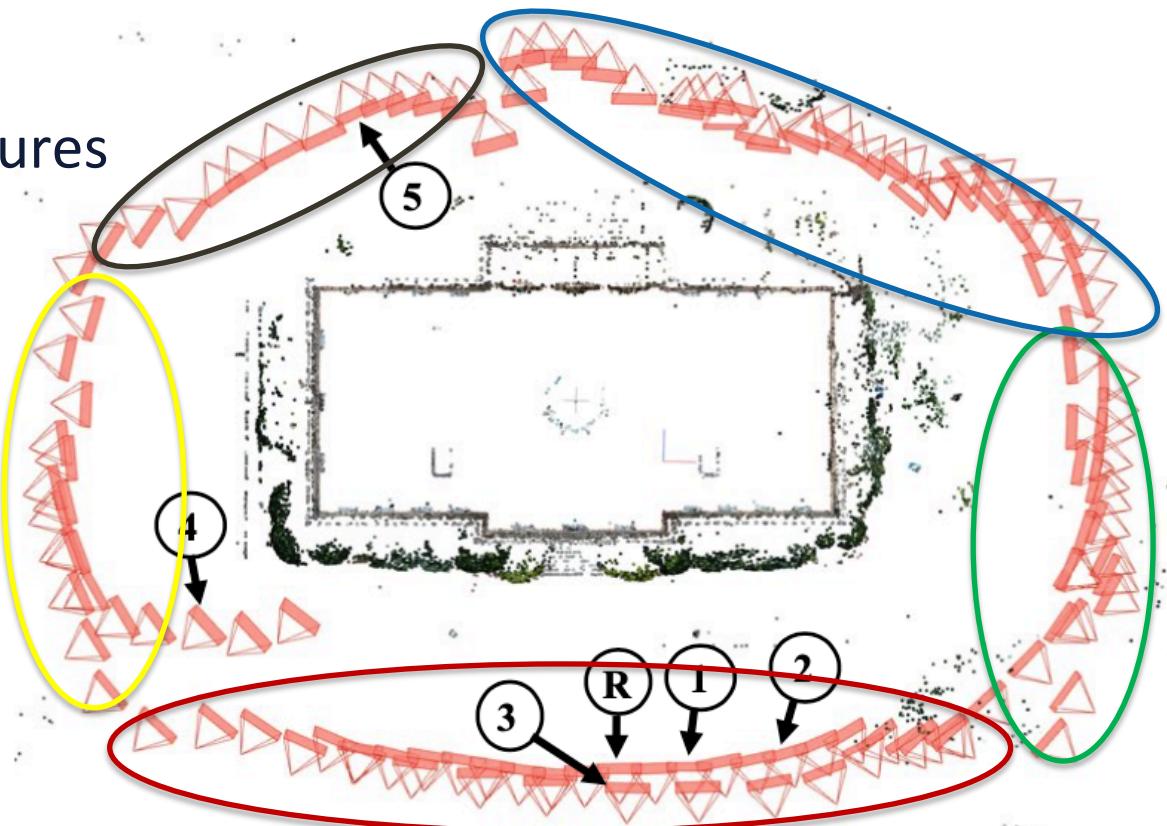
- Coarse visibility estimation via pose clustering



View Selection

- Coarse visibility estimation via pose clustering

- Camera Proximity
- Shared Sparse Features



View Selection

- **Fine-scale visibility estimation**
- *Good candidate source image ?*

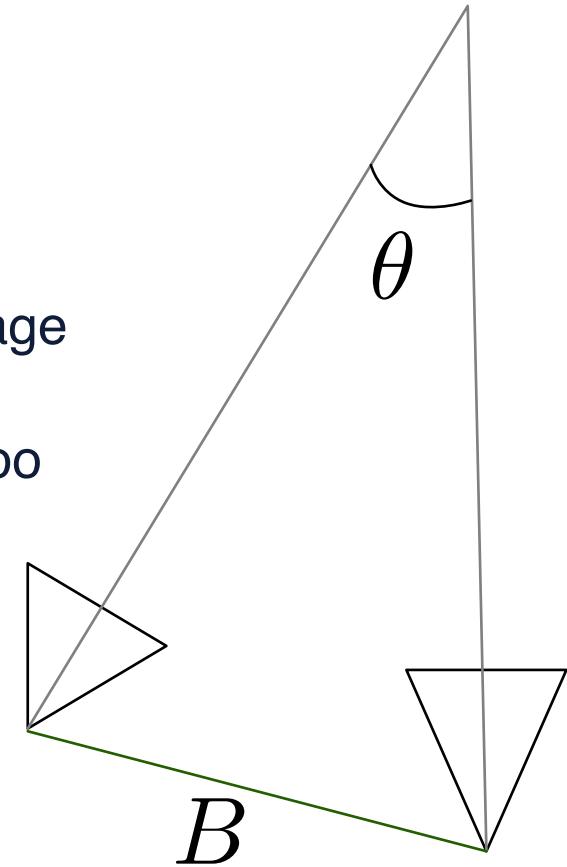
View Selection

- **Fine-scale visibility estimation**
 - *Good candidate source image ?*
 - Global
 - -> a **similar viewing direction** as the target image
 - -> a **suitable baseline** neither too short to degenerate the reconstruction accuracy nor too long to have less common coverage of the scene.

View Selection

- **Fine-scale visibility estimation**

- *Good candidate source image ?*
- Global
- -> a **similar viewing direction** as the target image
- -> a **suitable baseline** neither too short to degenerate the reconstruction accuracy nor too long to have less common coverage of the scene.
- $5^\circ < \theta < 60^\circ$
- $0.05d \leq B \leq 2d$



Pixel-Level View Selection

Reference



Source



Credit: E. Dunn

Pixel-Level View Selection

Reference



Source



Credit: E. Dunn

Pixel-Level View Selection

Reference



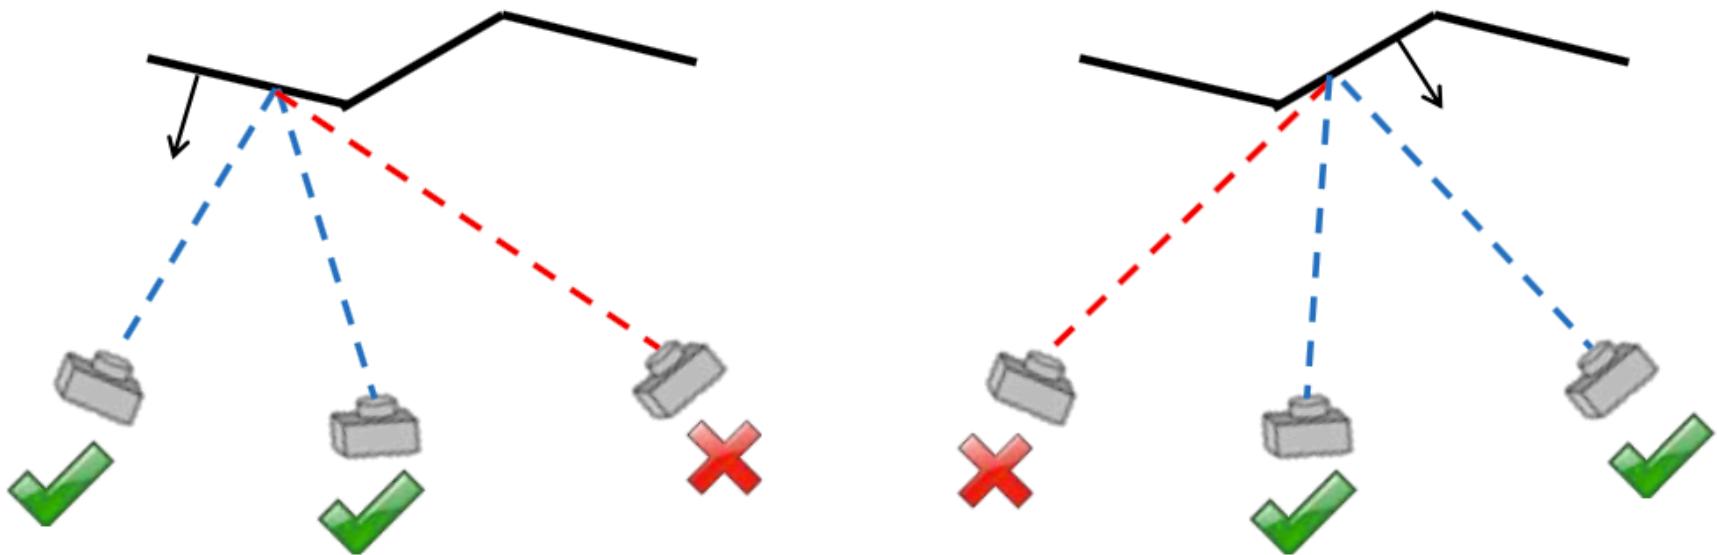
Source



Credit: E. Dunn

Image Selection vs Depth Estimation

- visibility requires scene structure and scene structure requires visibility
- **This is a chicken-and-egg problem**



Credit: S. Shen

Joint Pixel-Level View Selection and Depthmap Estimation

- Maximum likelihood estimation (MLE):
- Likelihood function $P(X, Z, \theta, N)$:
- Images:
 - $X = \{X^{ref}, X^{src}\}, X^{src} = \{X^m | m = 1 \dots M\}$
- Depth:
 - $\theta = \{\theta_l | l = 1 \dots L\}$
- Normal:
 - $N = \{n_l | l = 1 \dots L\}$
- Occlusion indicators:
 - $Z = \{Z_l^m | l = 1 \dots L, m = 1 \dots M\}, Z_l^m \in \{0, 1\}$

E. Zheng, et. al. “Patchmatch based joint view selection and depthmap estimation”, CVPR 2014

J. L. Schönberger, et. al. “Pixelwise View Selection for Unstructured Multi-View Stereo”, ECCV 2016

Joint Pixel-Level View Selection and Depthmap Estimation

- Maximum likelihood estimation (MLE):

- Likelihood function $P(X, Z, \theta, N)$:

$$\prod_{l=1}^L \prod_{m=1}^M [P(Z_{l,t}^m | Z_{l-1,t}^m, Z_{l,t-1}^m) P(X_l^m | Z_l^m, \theta_l, n_l) P(\theta_l, n_l | \theta_l^m, n_l^m)]$$

- Images:

- $X = \{X^{ref}, X^{src}\}, X^{src} = \{X^m | m = 1 \dots M\}$

- Depth:

- $\theta = \{\theta_l | l = 1 \dots L\}$

- Normal:

- $N = \{n_l | l = 1 \dots L\}$

- Occlusion indicators:

- $Z = \{Z_l^m | l = 1 \dots L, m = 1 \dots M\}, Z_l^m \in \{0, 1\}$

Joint Pixel-Level View Selection and Depthmap Estimation

- Maximum likelihood estimation (MLE):
- Likelihood function $P(X, Z, \theta, N)$:

$$\prod_{l=1}^L \prod_{m=1}^M [P(Z_{l,t}^m | Z_{l-1,t}^m, Z_{l,t-1}^m) P(X_l^m | Z_l^m, \theta_l, n_l) P(\theta_l, n_l | \theta_l^m, n_l^m)]$$


Photometric prior

If $Z_l^m = 1$,

$$P(X_l^m | Z_l^m, \theta_l, n_l) \propto \rho_l^m(\theta_l, n_l) \quad (\text{color similarity})$$

If $Z_l^m = 0$,

$$P(X_l^m | Z_l^m, \theta_l, n_l) = \text{uniform distribution}$$

Joint Pixel-Level View Selection and Depthmap Estimation

- Maximum likelihood estimation (MLE):
- Likelihood function $P(X, Z, \theta, N)$:

$$\prod_{l=1}^L \prod_{m=1}^M [P(Z_{l,t}^m | Z_{l-1,t}^m, Z_{l,t-1}^m) P(X_l^m | Z_l^m, \theta_l, n_l) P(\theta_l, n_l | \theta_l^m, n_l^m)]$$

Spatial-temporary smoothness

$$P(Z_{l,t}^m | Z_{l-1,t}^m, Z_{l,t-1}^m) = P(Z_{l,t}^m | Z_{l-1,t}^m) P(Z_{l,t}^m | Z_{l,t-1}^m).$$

$$P(Z_l^m | Z_{l-1}^m) = \begin{pmatrix} \gamma & 1-\gamma \\ 1-\gamma & \gamma \end{pmatrix}.$$

$$P(Z_{l,t}^m | Z_{l,t-1}^m) = \begin{pmatrix} \lambda_t & 1-\lambda_t \\ 1-\lambda_t & \lambda_t \end{pmatrix}$$

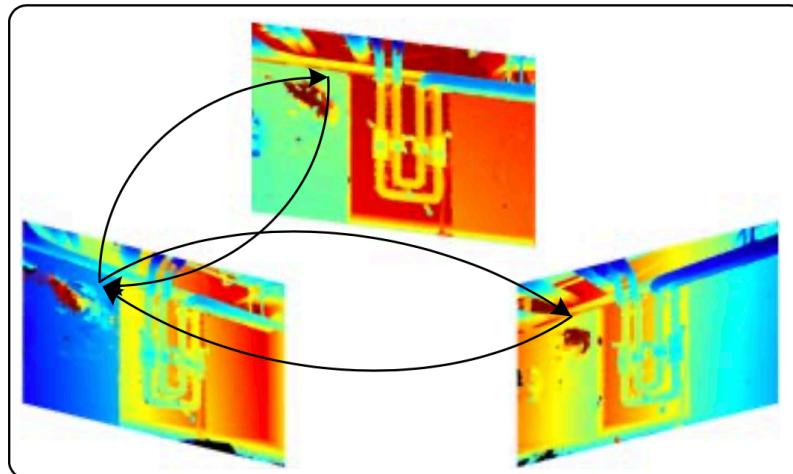
Joint Pixel-Level View Selection and Depthmap Estimation

- Maximum likelihood estimation (MLE):
- Likelihood function $P(X, Z, \theta, N)$:

$$\prod_{l=1}^L \prod_{m=1}^M [P(Z_{l,t}^m | Z_{l-1,t}^m, Z_{l,t-1}^m) P(X_l^m | Z_l^m, \theta_l, n_l) P(\theta_l, n_l | \theta_l^m, n_l^m)]$$

geometric consistency

Forward-backward
reprojection error



E. Zheng, et. al. “Patchmatch based joint view selection and depthmap estimation”, CVPR 2014

J. L. Schönberger, et. al. “Pixelwise View Selection for Unstructured Multi-View Stereo”, ECCV 2016

Joint Pixel-Level View Selection and Depthmap Estimation

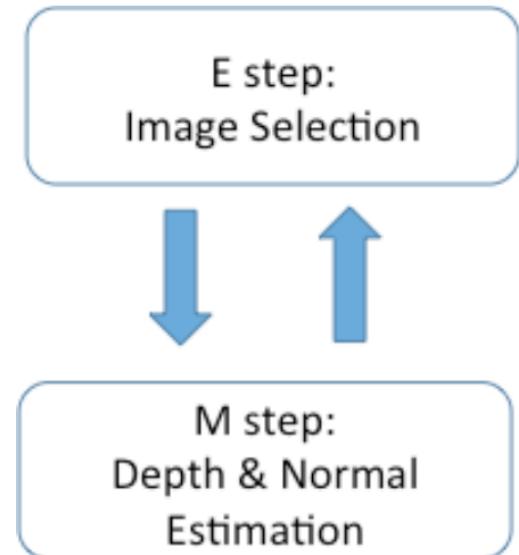
- Joint likelihood Estimation:

$$P(X, Z, \theta, N)$$

The diagram illustrates the joint likelihood function $P(X, Z, \theta, N)$. A bracket on the left side of the equation branches out to four arrows pointing to the right, each labeled with a variable: Normals, Depth, Occlusion, and Images.

- Generalized Expectation Maximization

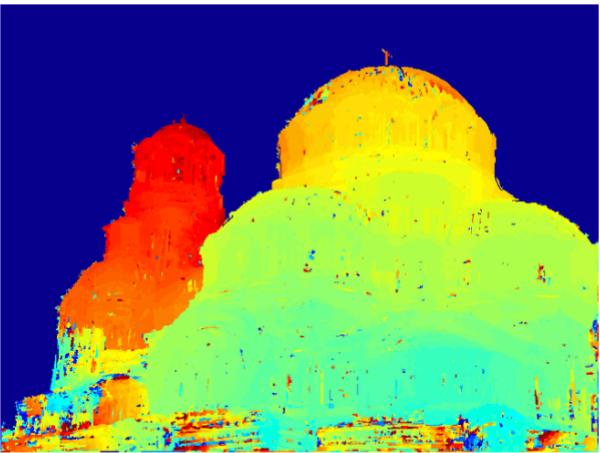
- E-Step
 - Infer Z using variational inference
- M-Step
 - Infer θ, N using PatchMatch sampling



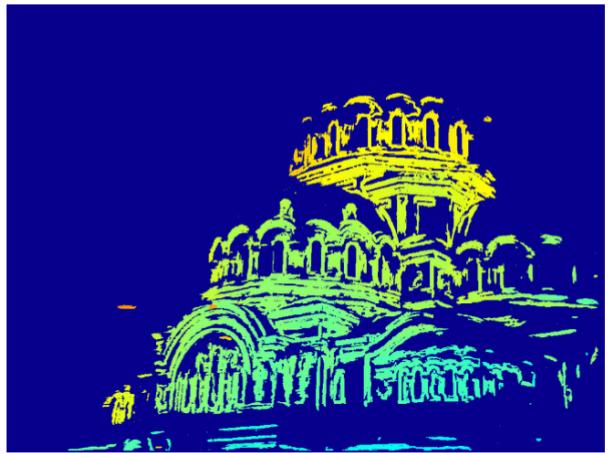
E. Zheng, et. al. “Patchmatch based joint view selection and depthmap estimation”, CVPR 2014

J. L. Schönberger, et. al. “Pixelwise View Selection for Unstructured Multi-View Stereo”, ECCV 2016

Robustness of Pixel-Level Selection

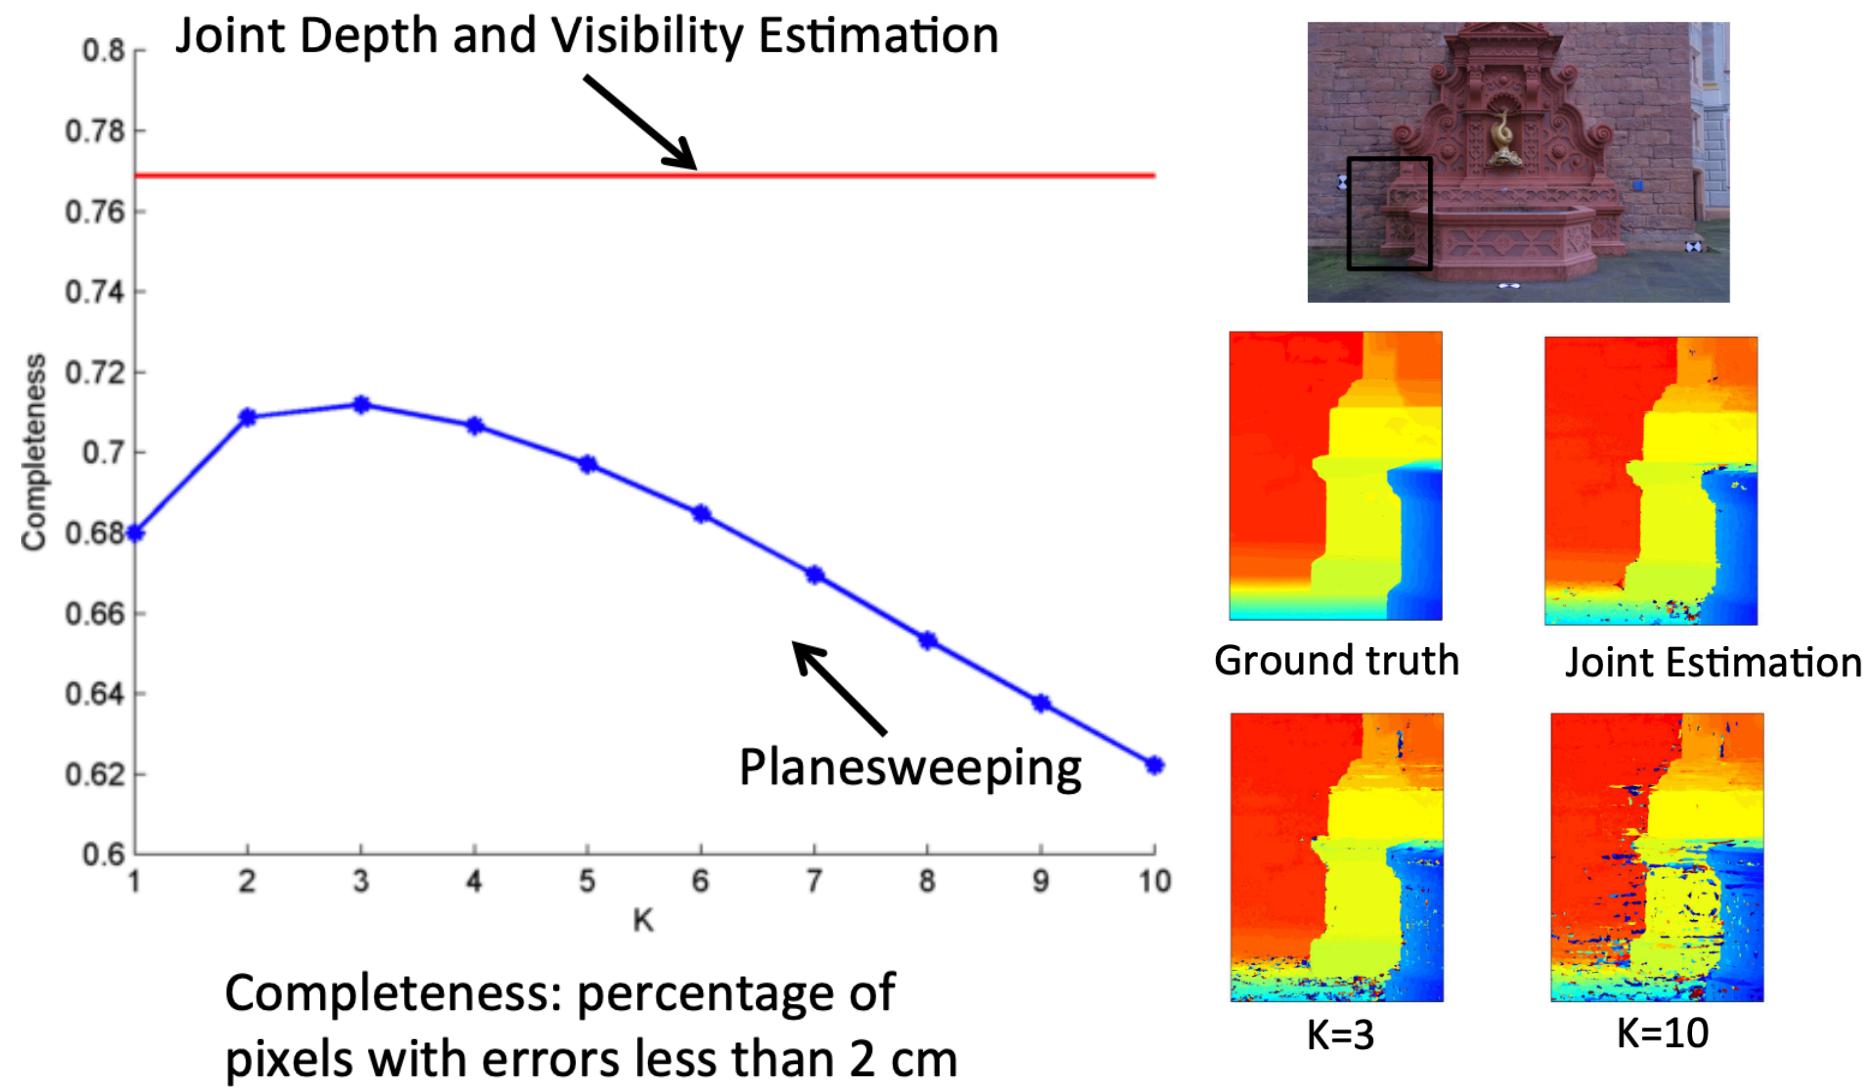


Pixel-level



Baseline

Robustness of Pixel-Level Selection



Summary

- **PatchMatch**
 - A randomized algorithm for rapidly finding correspondences between image patches
 - Step
 - 1: Initialization
 - 2: Propagation
 - 3: Random Search
- **PatchMatch Stereo**
 - Finding a “good” slanted support plane at each pixel.
 - Difference from vanilla PatchMatch
 - (offset) -> (depth, normal)
- **View Selection**
 - Coarse visibility estimation
 - Fine-scale visibility estimation
 - Joint Pixel-Level View Selection and Depthmap Estimation
 - Pixel-Level occlusion indicator
 - chicken-and-egg -> Generalized Expectation maximization

Thanks